

- This presentation presents a broad overview of methods for interpreting interactions in logistic regression.
- The presentation is not about Stata. It uses Stata, but you gotta use something.
- The methods shown are somewhat stat package independent. However, they can be easier or more difficult to implement depending on the stat package.
- The presentation is not a step-by-step how-to manual that shows all of the code that was used to produce the results shown.
- Each of the models used in the examples will have two research variables that are interacted and one continuous covariate (**cv1**) that is not part of the interaction.

## Some Definitions

### Odds

Showing that odds are ratios.

$\text{odds} = p / (1 - p)$
-----------------------------

### Log Odds

Natural log of the odds, also known as a logit.

$\log \text{ odds} = \text{logit} = \log(p / (1 - p))$
--

### Odds Ratio

Showing that odds ratios are actually ratios of ratios.

$\text{odds\_ratio} = \frac{\text{odds1}}{\text{odds2}} = \frac{p1 / (1 - p1)}{p2 / (1 - p2)}$
--

### Computing Odds Ratio from Logistic Regression Coefficient

$\text{odds\_ratio} = \exp(b)$
--------------------------------

### Computing Probability from Logistic Regression Coefficients

$\text{probability} = \exp(Xb) / (1 + \exp(Xb))$
--

Where **Xb** is the linear predictor.

## About Logistic Regression

Logistic regression fits a maximum likelihood logit model. The model estimates conditional means in terms of logits (log odds). The logit model is a linear model in the log odds metric. Logistic regression results can be displayed as odds ratios or as probabilities. Probabilities are a nonlinear transformation of the log odds results.


In general, linear models have a number of advantages over nonlinear models and are easier to work with. For example, in linear models the slopes and/or differences in means do not change for differing values of a covariate. This is not necessarily the case for nonlinear models. The problem in logistic regression is that, even though the model is linear in log odds, many researchers feel that log odds are not a natural metric and are not easily interpreted.

Probability is a much more natural metric. However, the logit model is not linear when working in the probability metric. Thus, the predicted probabilities change as the values of a covariate change. In fact, the estimated probabilities depend on all variables in the model not just the variables in the interaction.

So what is a linear model? A linear model is linear in the betas (coefficients). By extension, a nonlinear model must be nonlinear in the betas. Below are three example of linear and nonlinear models.


First, is an example of a linear model and its graph.

$$\hat{Y} = \beta_0 + \beta_1 * X$$




Next we have an example of a nonlinear model and its graph. In this case its an exponential growth model.

$$\hat{Y} = \beta_0 * \beta_1^X$$




Lastly we have another nonlinear model. This one shows the nonlinear transformation of log odds to probabilities.

$$\hat{Y} = \frac{e^{(\beta_0 + \beta_1 * X)}}{(1 + e^{(\beta_0 + \beta_1 * X)})}$$



Logistic Regression Transformations

This is an attempt to show the different types of transformations that can occur with logistic regression models.



## Logistic interactions are a complex concept

Common wisdom suggests that interactions involves exploring differences in differences. If the differences are not different then there is no interaction. But in logistic regression interaction is a more complex concept. Researchers need to decide on how to conceptualize the interaction. Is the interaction to be conceptualized in terms of log odds (logits) or odds ratios or probability? This decision can make a big difference. An interaction that is significant in log odds may not be significant in terms of difference in differences for probability. Or *vice versa*.

## Model 1: categorical by categorical interaction

### Log odds metric — categorical by categorical interaction

Variables **f** and **h** are binary predictors, while **cv1** is a continuous covariate. The **nolog** option suppresses the display of the iteration log; it is used here simply to minimize the quantity of output.

logit y01 f##h cv1, nolog						
Logistic regression			Number of obs		=	200
			LR chi2(4)		=	106.10
			Prob > chi2		=	0.0000
Log likelihood = -78.74193			Pseudo R2		=	0.4025
-----						
y01		Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
-----+-----						
1.f		2.996118	.7521524	3.98	0.000	1.521926 4.470309
1.h		2.390911	.6608498	3.62	0.000	1.09567 3.686153
f#h						
1 1		-2.047755	.8807989	-2.32	0.020	-3.774089 -.3214213
cv1		.196476	.0328518	5.98	0.000	.1320876 .2608644
_cons		-11.86075	1.895828	-6.26	0.000	-15.5765 -8.144991
-----						

The interaction term is clearly significant. We could manually compute the expected logits for each of the four cells in the model.



```

/* difference 1 at f = 0 */

lincom 0.f#0.h - 0.f#1.h

( 1)  [y01]0bn.f#0bn.h - [y01]0bn.f#1.h = 0

```

y01	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
(1)	-2.390911	.6608498	-3.62	0.000	-3.686153	-1.09567

```

/* difference 2 at f = 1 */

lincom 1.f#0.h - 1.f#1.h

( 1)  [y01]1.f#0bn.h - [y01]1.f#1.h = 0

```

y01	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
(1)	-.3431562	.5507722	-0.62	0.533	-1.42265	.7363375

Difference 1 suggests that **h0** is significantly different from **h1** at **f** = 0, While difference 2 does not show a significant difference at **f** = 1. These are tests of simple main effects just like we would do in OLS (ordinary least squares) regression. We will finish up this section by looking at the difference in differences.

```
/* difference in differences */

lincom (0.f#0.h - 0.f#1.h)-(1.f#0.h - 1.f#1.h)

( 1)  [y01]0bn.f#0bn.h - [y01]0bn.f#1.h - [y01]1.f#0bn.h + [y01]1.f#1.h = 0
```

-----+-----						
y01	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----						
(1)	-2.047755	.8807989	-2.32	0.020	-3.774089	-.3214213
-----+-----						

The difference in differences is, of course, just another name for the interaction. For the log odds model the differences and the difference in differences are the same regardless of the value of the covariate. This constancy across different values of the covariate is one of the properties of linear models.

### Odds ratio metric — categorical by categorical interaction

Let’s look at a table of logistic regression coefficients along with the exponentiated coefficients, which some people call odds ratios.

source	coefficient	exp(coef)	type of exp(coef)
f	2.996118	20.007716	odds ratio
h	2.390911	10.92345	odds ratio
f#h	-2.047755	0.1290242	ratio of odds ratios
cv1	0.196476	1.217106	odds ratio
_cons	-11.86075	7.062e-06	baseline odds

Many people call all exponentiated logistic coefficients odds ratios. But as you can see from the table above, exponentiating the interaction is a ratio of ratios and the exponentiated constant is the baseline odds.

We can compute the odds ratios manually for each of the two levels of **f** from the values in the table above.

<b>odds ratio h1/h0 for f=0:   b[1.h]</b>	<b>= 10.92345</b>
<b>odds ratio h1/h0 for f=1:   b[1.h]*b[f#h] = 10.92345*.1290242</b>	<b>= 1.4093894</b>

Please note that the computation of the odds ratio for **f** =1 involves multiplying coefficients for the odds ratio model above which implies that odds ratio models are multiplicative rather than additive.

The **baseline** odds when **cv1** = zero is very small (7.06e-06) so for the remainder of of the computations we will estimate the odds while holding **cv1** at 50. The option **noatlegend** suppresses the display of the legend.

margins, over(f h) at(cv1=50) expression(exp(xb())) noatlegend							
Predictive margins				Number of obs		=	200
Model VCE		: OIM					
Expression		: exp(xb())					
over		: f h					
-----+-----							
		Delta-method					
		Margin	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
f h							
0 0		.1304264	.0734908	1.77	0.076	-.0136129	.2744657
0 1		1.424706	.515989	2.76	0.006	.4133857	2.436025
1 0		2.609533	1.136545	2.30	0.022	.3819457	4.837121
1 1		3.677847	1.311463	2.80	0.005	1.107427	6.248267
-----+-----							

The option **expression(exp(xb()))** insures that we are looking at results in the odds ratio metric. The **baseline** odds are now .1304264 which is reasonable. We will compute the odds ratio for each level of **f**.

<b>odds ratio 1 at f=0: 1.424706/.1304264 = 10.923446</b>
<b>odds ratio 2 at f=1: 3.677847/2.609533 = 1.4093889</b>

So when **f** = 0 the odds of the outcome being one are 10.92 times greater for **h1** then for **h0**. For **f** = 1 the ratio of the two odds is only 1.41. These odds ratios are the same as we computed manually earlier.

We can also compute the ratio of odds ratios and show that it reproduces the estimate for the interaction.

<b>ratio of odds ratios: (3.677847/2.609533)/(1.424706/.1304264) = .1290242</b>
---

The one nice thing that we can say about working in odds ratio metric is the odds ratios remain the same regardless of where we hold the covariate constant.

## Probability metric — categorical by categorical interaction

We will begin by rerunning our logistic regression model to refresh our memories on the coefficients.

logit y01 f##h cv1, nolog							
Logistic regression				Number of obs		=	200
				LR chi2(4)		=	106.10
				Prob > chi2		=	0.0000
Log likelihood = -78.74193				Pseudo R2		=	0.4025
-----							
y01		Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
1.f		2.996118	.7521524	3.98	0.000	1.521926	4.470309
1.h		2.390911	.6608498	3.62	0.000	1.09567	3.686153
f#h							
1 1		-2.047755	.8807989	-2.32	0.020	-3.774089	-.3214213
cv1		.196476	.0328518	5.98	0.000	.1320876	.2608644
_cons		-11.86075	1.895828	-6.26	0.000	-15.5765	-8.144991
-----							

Let’s manually compute the probability of the outcome being one for the **f** = 0, **h** = 0 cell when **cv1** is held at 50.

<b>Xb = b[_cons] + 0*b[1.f] + 0*b[1.h] + 0*b{f#h} + 50*b[cv1]</b>
<b>= -11.86075 + 0*2.996118 + 0*2.390911 + 0*-2.047755 + 50*.196476 = -2.03695</b>
<b>probability = exp(Xb)/(1+exp(Xb)) = exp(-2.03695)/(1+exp(-2.03695)) = .11537767</b>

We could repeat this for each of the other three cells but instead we we will obtain the expected probabilities for each cell while holding the covariate at 50 using the **margins** command.

margins f#h, at(cv1=50)

Adjusted predictions

Number of obs = 200

Model VCE : OIM

Expression : Pr(y01), predict()

at : cv1 = 50

		Delta-method				
		Margin	Std. Err.	z	P> z	[95% Conf. Interval]
f#h						
0 0		.115378	.0575106	2.01	0.045	.0026592 .2280968
0 1		.5875788	.0877652	6.69	0.000	.4155621 .7595955
1 0		.7229559	.0872338	8.29	0.000	.5519808 .8939309
1 1		.7862264	.0599327	13.12	0.000	.6687605 .9036924

Here are the same results displayed as a table.

	h=0	h=1
f=0	.115378	.5875788
f=1	.7229559	.7862264

We would like to look at the differences in **h** for each level of **f**.

<b>h1 - h0 at f = 0: .5875788 - .115378 = .4722008</b>		
<b>h1 - h0 at f = 1: .7862264 - .7229559 = .0632706</b>		

We can also do this with a slight variation of the **margins** command and get estimates of the differences in probability along with standard errors and confidence intervals.



margins f, dydx(h) at(cv1=50) post							
Conditional marginal effects				Number of obs		=	200
Model VCE		: OIM					
Expression		: Pr(y01), predict()					
dy/dx w.r.t.		: 0.h 1.h					
at		: cv1		=	50		
-----							
		Delta-method					
		dy/dx	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
1.h							
	f						
	0	.4722008	.1035128	4.56	0.000	.2693195	.675082
	1	.0632706	.1036697	0.61	0.542	-.1399183	.2664595
-----							
Note: dy/dx for factor levels is the discrete change from the base level.							

These two differences are the probability analogs to the simple main effects from the log odds model. So, when the covariate is held at 50 there is a significant difference in **h** at **f** = 0 but not at **f** = 1.

Next, we will use **lincom** to compute the difference in differences when **cv1** is held at 50.

lincom [1.h]0.f-[1.h]1.f

( 1) [1.h]0bn.f - [1.h]1.f = 0

	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
(1)	.4089302	.1482533	2.76	0.006	.118359	.6995014

The p-value here is different form the p-value from the original logit model because in the probability metric the values of the covariate matter.

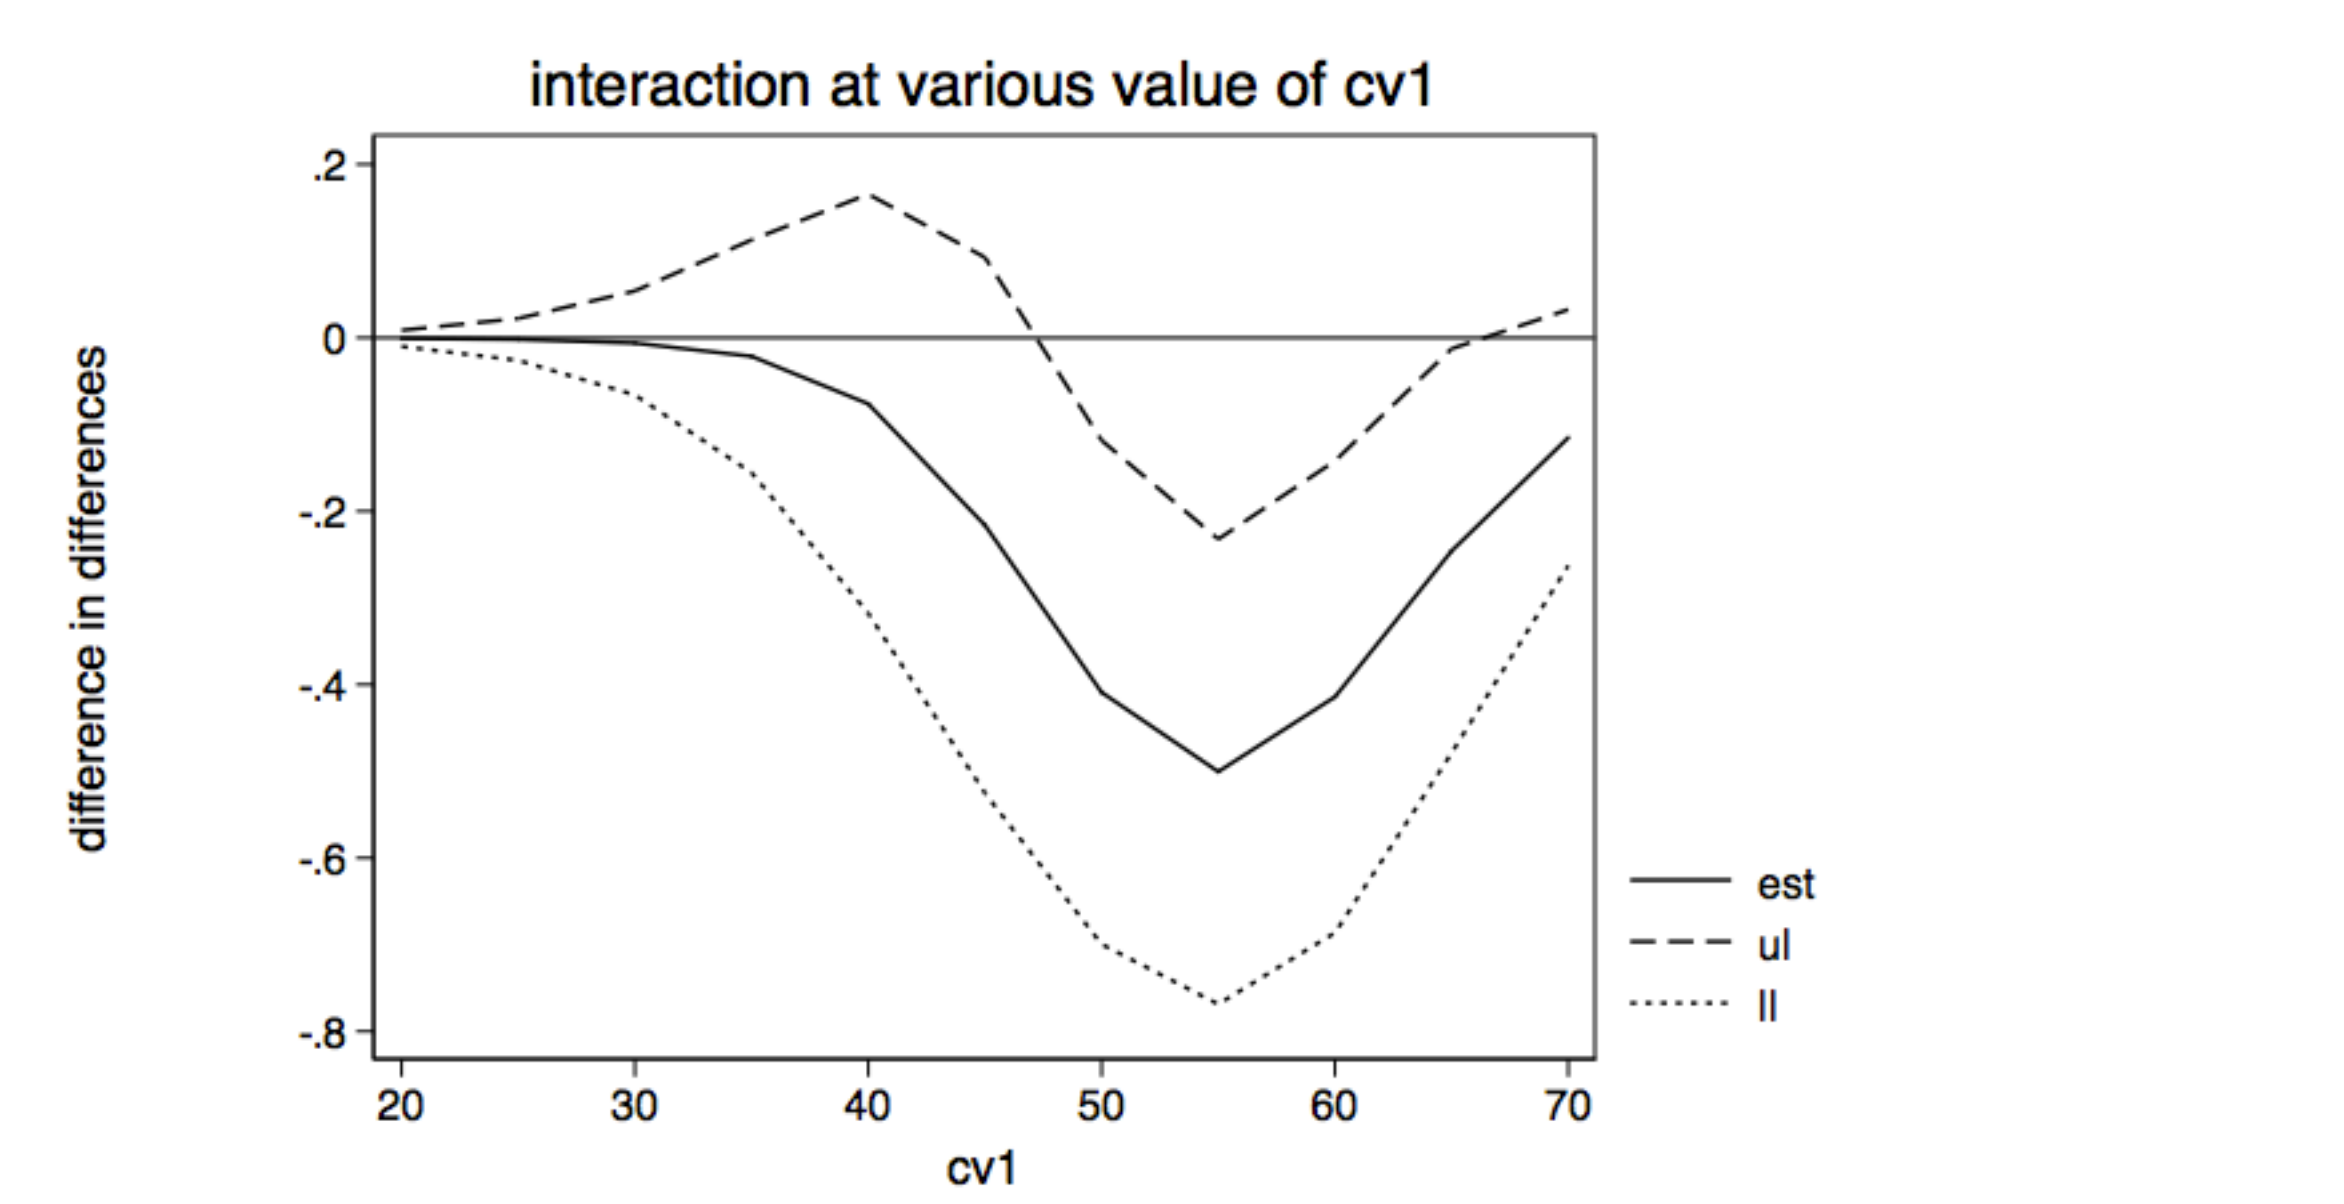
If we repeat the above process for values of **cv1** from 20 to 70, we can produce a table of simple main effects and a graph of the difference in differences.

Table of Simple Main Effects for h at Two Levels of f for Various Values of cv1

		Delta-method					
		dy/dx	Std. Err.	z	P> z	[95% Conf. Interval]	
cv1 f							
20	0	.0035507	.0038256	0.93	0.353	-.0039472	.0110487
20	1	.002893	.0057719	0.50	0.616	-.0084197	.0142058
30	0	.0246805	.0188412	1.31	0.190	-.0122475	.0616086
30	1	.0186252	.0331697	0.56	0.574	-.0463863	.0836367
40	0	.1485222	.0656193	2.26	0.024	.0199107	.2771337
40	1	.0723494	.1167547	0.62	0.535	-.1564856	.3011843
50	0	.4722008	.1035128	4.56	0.000	.2693195	.675082
50	1	.0632706	.1036697	0.61	0.542	-.1399183	.2664595
60	0	.4284548	.137549	3.11	0.002	.1588636	.6980459
60	1	.0142654	.0255894	0.56	0.577	-.0358888	.0644197
70	0	.1173445	.076704	1.53	0.126	-.0329926	.2676816
70	1	.0021597	.0042758	0.51	0.613	-.0062207	.0105402

Table of Difference in Differences for Various Values of cv1

		Delta-method					
		dy/dx	Std. Err.	z	P> z	[95% Conf. Interval]	
cv1							
20		.0006577	.0047463	0.14	0.890	-.0086449	.0099603
30		.0060553	.0306291	0.20	0.843	-.0539766	.0660872
40		.0761728	.1233778	0.62	0.537	-.1656432	.3179889
50		.4089302	.1482533	2.76	0.006	.118359	.6995014
60		.4141893	.1388141	2.98	0.003	.1421186	.68626
70		.1151848	.0753487	1.53	0.126	-.0324959	.2628654



Clearly, the value of the covariate makes a huge difference in whether or not the simple main effects or the interactions are statistically significant when working in the probability metric.

## Model 1a: Categorical by categorical interaction?

But wait, what if the model does not contain an interaction term? Consider the following model.

logit y01 i.f i.h cv1						
Logistic regression			Number of obs		=	200
			LR chi2(3)		=	100.26
			Prob > chi2		=	0.0000
Log likelihood = -81.6618			Pseudo R2		=	0.3804
-----						
y01		Coef.	Std. Err.	z	P> z	[95% Conf. Interval]
-----+-----						
1.f		1.65172	.4229992	3.90	0.000	.8226566 2.480783
1.h		1.256555	.4009757	3.13	0.002	.4706575 2.042453
cv1		.1806214	.0304036	5.94	0.000	.1210314 .2402113
_cons		-10.26943	1.622842	-6.33	0.000	-13.45015 -7.088723
-----						

We will manually compute the expected log odds for each of the four cells of the model.

f	h
cell 0 0	b[_cons]
	= -10.26943
cell 1 0	b[_cons] + b[1.f]
	= -10.26943 + 1.65172 = -8.61771
cell 0 1	b[_cons] + b[1.h]
	= -10.26943 + 1.256555 = -9.012875
cell 1 1	b[_cons] + b[1.f] + b[1.h]
	= -10.26943 + 1.65172 + 1.256555 = -7.361155

Next we will compute the differences for **f**=0 and **f**=1.

difference 1 at f = 0: -10.26943 - -8.6177 = -1.65173

difference 2 at f = 1: -9.012875 - -7.361155 = -1.65172

They are identical to within rounding error, showing that there is no interaction effect in the log odds model.


Next we will compute the expected probabilities for **cv1** held at 50 along with the difference in differences.

margins, over(f h) at(cv1=50) post							
Predictive margins				Number of obs		=	200
-----							
		Delta-method					
		Margin	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
f#h							
0 0		.2247204	.0670438	3.35	0.001	.0933171	.3561238
0 1		.5045471	.0798579	6.32	0.000	.3480285	.6610657
1 0		.6018917	.0866773	6.94	0.000	.4320073	.7717761
1 1		.8415636	.0455686	18.47	0.000	.7522509	.9308764
-----							
lincom (_b[0.f#1.h]-_b[0.f#0.h])-(_b[1.f#1.h]-_b[1.f#0.h])							
( 1) - 0bn.f#0bn.h + 0bn.f#1.h + 1.f#0bn.h - 1.f#1.h = 0							
-----							
		Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
(1)		.0401547	.0364121	1.10	0.270	-.0312117	.111521
-----							

The difference in differences is not very large. Let's try in again, this time holding **cv1** at 60.

margins, over(f h) at(cv1=60) post							
Predictive margins				Number of obs		=	200
-----							
		Delta-method					
		Margin	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
f#h							
0 0		.6382663	.1046912	6.10	0.000	.4330753	.8434572
0 1		.8610935	.0455552	18.90	0.000	.7718069	.9503802
1 0		.9019929	.0470231	19.18	0.000	.8098294	.9941565
1 1		.9700007	.0146765	66.09	0.000	.9412353	.998766
-----							
lincom (_b[0.f#1.h]-_b[0.f#0.h])-(_b[1.f#1.h]-_b[1.f#0.h])							
( 1) - 0bn.f#0bn.h + 0bn.f#1.h + 1.f#0bn.h - 1.f#1.h = 0							
-----							
		Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
(1)		.1548195	.0634635	2.44	0.015	.0304334	.2792057
-----							

This time the difference in differences is much larger. Let’s make a graph similar to the one we did for the model with the interaction included.



We see that, even without an interaction term in the model, the differences in differences (interactions?) can vary widely from negative to positive depending on the value of the covariate.

This leads us to the “Quote of the Day.”

## Quote of the day

Departures from additivity imply the presence of interaction types, but additivity does **not** imply the absence of interaction types.

*Greenland & Rothman, 1998*

## Model 2: Categorical by continuous interaction

### Log odds metric — categorical by continuous interaction

The dataset for the categorical by continuous interaction has one binary predictor (**f**), one continuous predictor (**s**) and a continuous covariate (**cv1**). Let’s take a look at the logistic regression model.

logit y f##c.s cv1							
Logistic regression				Number of obs		=	200
				LR chi2(4)		=	114.41
				Prob > chi2		=	0.0000
Log likelihood = -74.587842				Pseudo R2		=	0.4340
-----							
y		Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
1.f		9.983662	3.05269	3.27	0.001	4.0005	15.96682
s		.1750686	.0470033	3.72	0.000	.0829438	.2671933
f#c.s							
1		-.1595233	.0570352	-2.80	0.005	-.2713103	-.0477363
cv1		.1877164	.0347888	5.40	0.000	.1195316	.2559013
_cons		-19.00557	3.371064	-5.64	0.000	-25.61273	-12.39841
-----							


The interaction term is significant indicating the the slopes for **y** on **s** are significantly different for each level of **f**. We can compute the slopes and intercepts manually as shown below.

<b>slope for f=0:    b[s] = .1750686</b>	
<b>slope for f=1:    b[s] + b[f#c.s] = .1750686 - .1595233 = .0155453</b>	
<b>intercept for f=0:    _cons = -19.00557</b>	
<b>intercept for f=1:    _cons + b[1.f]= -19.00557 + 9.983662 = -9.021909</b>	

Here are our two logistic regression equations in the log odds metric.

$-19.00557 + .1750686*s + 0*cv1$
$-9.021909 + .0155453*s + 0*cv1$

Now we can graph these two regression lines to get an idea of what is going on.



Because the logistic regress model is linear in log odds, the predicted slopes do not change with differing values of the covariate.

## Probability metric — categorical by continuous interaction

We'll begin by rerunning the logistic regression model.

logit y f##c.s cv1

Logistic regression

Number of obs = 200

LR chi2(4) = 114.41

Prob > chi2 = 0.0000

Log likelihood = -74.587842

Pseudo R2 = 0.4340

y	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
1.f	9.983662	3.05269	3.27	0.001	4.0005	15.96682
s	.1750686	.0470033	3.72	0.000	.0829438	.2671933
f#c.s						
1	-.1595233	.0570352	-2.80	0.005	-.2713103	-.0477363
cv1	.1877164	.0347888	5.40	0.000	.1195316	.2559013
_cons	-19.00557	3.371064	-5.64	0.000	-25.61273	-12.39841

If we were so inclined we could compute all of the probabilities of interest using the basic probability formula.

$\text{Prob} = \exp(\text{Xb}) / (1 + \exp(\text{Xb}))$
---

Here’s an example of computing the probability when f = 0, s = 60, f#s = 0, and cv1 =40.

$\text{Xb0} = -19.00557 + 0*9.983662 + 60*.1750686 + 0*-.1595233 + 40*.1877164 = -.992798$
$\exp(\text{Xb0}) / (1 + \exp(\text{Xb0})) = \exp(-.992798) / (1 + \exp(-.992798)) = .27035977$

Now we will use f = 1, s = 60, f#s = 60, and cv1 =40.

$\text{Xb1} = -19.00557 + 1*9.983662 + 60*.1750686 + 60*-.1595233 + 40*.1877164 = -.580534$
$\exp(\text{Xb1}) / (1 + \exp(\text{Xb1})) = \exp(-.580534) / (1 + \exp(-.580534)) = .35880973$

We can also compute the difference in probabilities.



$$\frac{\exp(\mathbf{Xb1})}{(1+\exp(\mathbf{Xb1}))} - \frac{\exp(\mathbf{Xb0})}{(1+\exp(\mathbf{Xb0}))} =$$
$$\frac{\exp(-.580534)}{(1+\exp(-.580534))} - \frac{\exp(-.992798)}{(1+\exp(-.992798))} = .08844995$$

If we use something like Stata’s **margins** command, we can get predicted probabilities along with standard errors and confidence intervals. Here is an example predicting the probability when **s** = 20 and **cv1** = 40.

margins f, at(s=20 cv1=40)							
Adjusted predictions				Number of obs		= 200	
Model VCE : OIM							
Expression : Pr(y), predict()							
-----+-----							
		Delta-method					
		Margin	Std. Err.	z	P> z	[95% Conf. Interval]	
-----+-----							
f							
0		.0003368	.0005779	0.58	0.560	-.0007958	.0014695
1		.2310582	.1500289	1.54	0.124	-.0629931	.5251095
-----+-----							


Now can repeat this for various values of s running from 20 to 70, producing the table below.

Table of Predicted Probabilities of f for Various Values of s Holding cv1 at 40


		Delta-method					
		Margin	Std. Err.	z	P> z	[95% Conf. Interval]	
s f							
20	0	.0003368	.0005779	0.58	0.560	-.0007958	.0014695
20	1	.2310582	.1500289	1.54	0.124	-.0629931	.5251095
25	0	.000808	.0012067	0.67	0.503	-.0015571	.003173
25	1	.2451555	.1320954	1.86	0.063	-.0137469	.5040578
30	0	.0019367	.0024706	0.78	0.433	-.0029056	.0067789
30	1	.2598222	.1136085	2.29	0.022	.0371536	.4824908
35	0	.0046348	.0049337	0.94	0.348	-.005035	.0143047
35	1	.2750467	.0959104	2.87	0.004	.0870657	.4630276
40	0	.0110505	.0095531	1.16	0.247	-.0076733	.0297743
40	1	.2908127	.081642	3.56	0.000	.1307973	.4508282
45	0	.0261139	.0178944	1.46	0.144	-.0089585	.0611863
45	1	.3070997	.0752299	4.08	0.000	.1596518	.4545475
50	0	.0604557	.0329478	1.83	0.067	-.0041208	.1250322
50	1	.3238822	.0808248	4.01	0.000	.1654685	.4822959
55	0	.1337569	.0622149	2.15	0.032	.0118178	.2556959
55	1	.3411303	.0980782	3.48	0.001	.1489005	.5333601
60	0	.2703596	.1168105	2.31	0.021	.0414151	.499304
60	1	.3588096	.1233704	2.91	0.004	.117008	.6006111
65	0	.4706697	.180248	2.61	0.009	.11739	.8239493
65	1	.3768809	.1535731	2.45	0.014	.0758831	.6778787
70	0	.6808947	.1951477	3.49	0.000	.2984123	1.063377
70	1	.3953013	.1867987	2.12	0.034	.0291827	.7614199

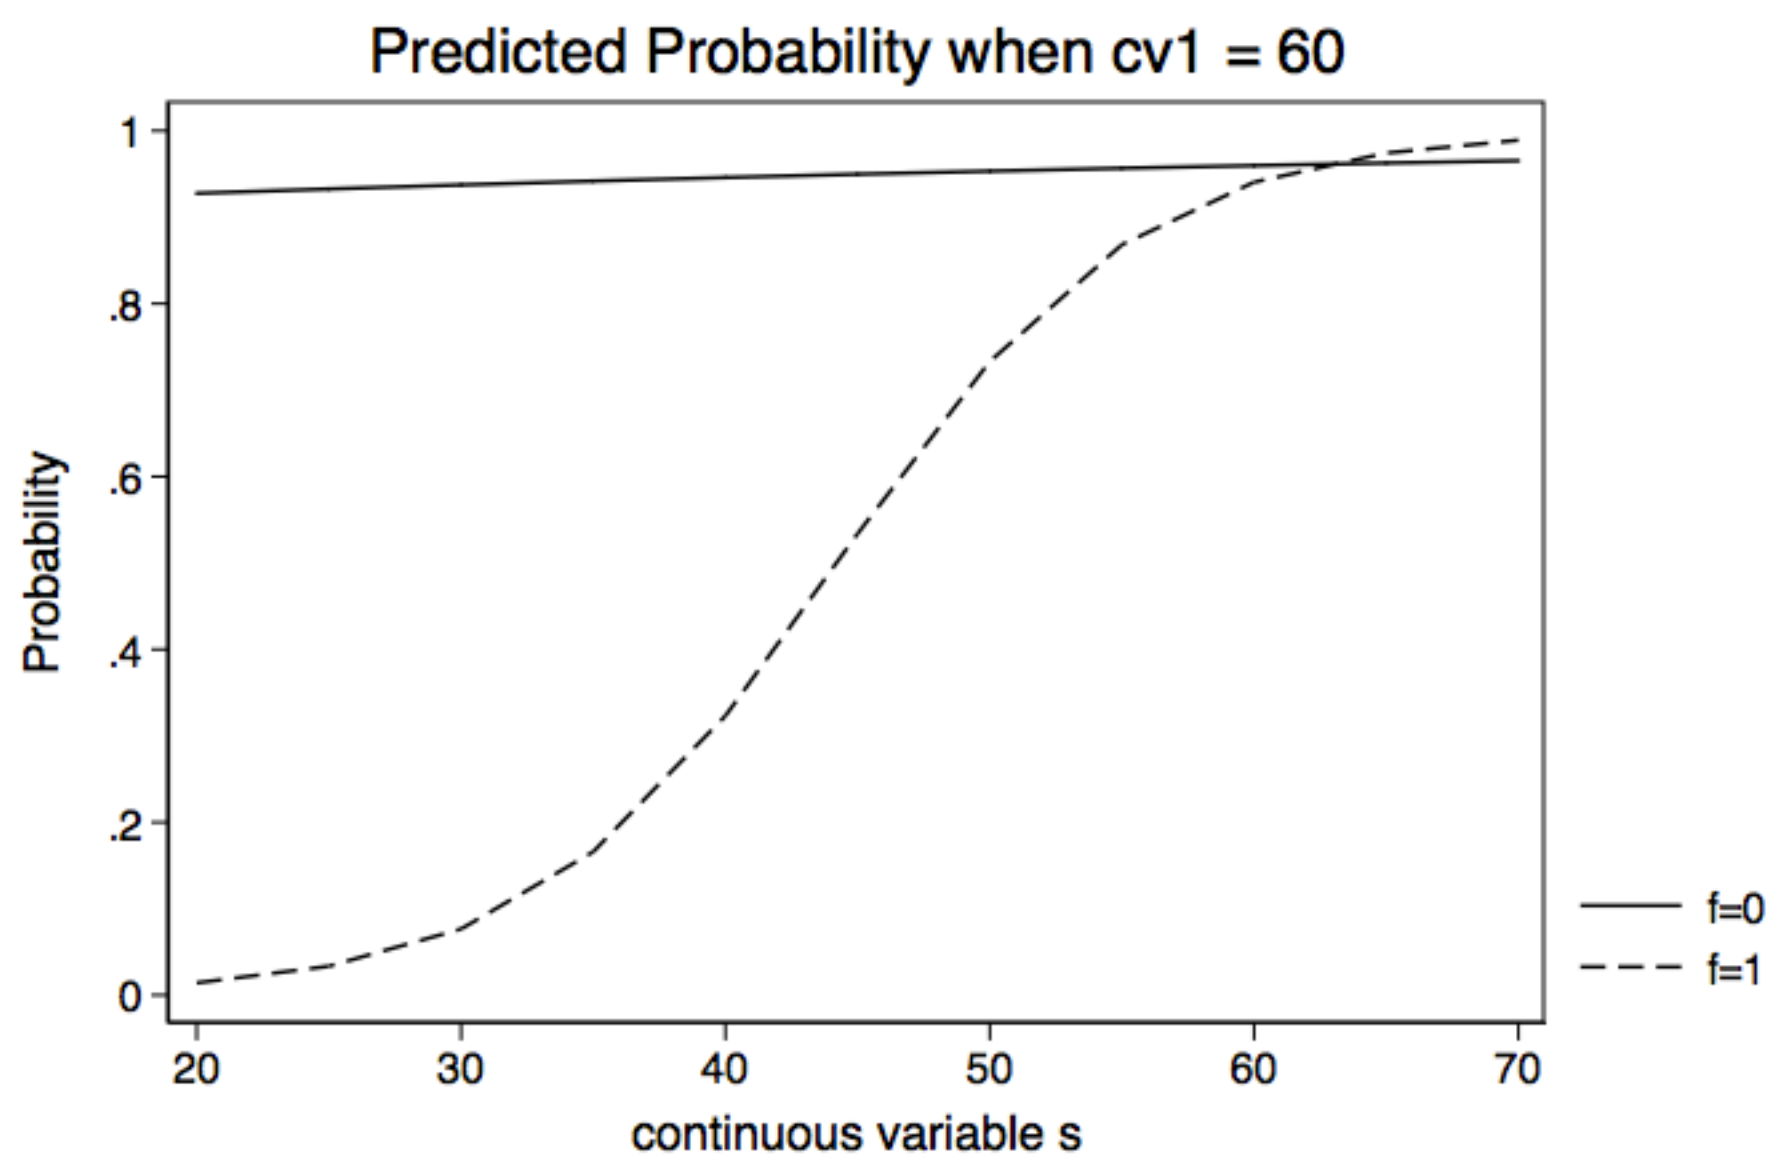
We will repeat this holding **cv1** at 50 and then 60. We will then plot the probabilities for each of the three values of **cv1**.

Predicted Probability when  $cv1 = 40$



Predicted Probability when  $cv1 = 50$





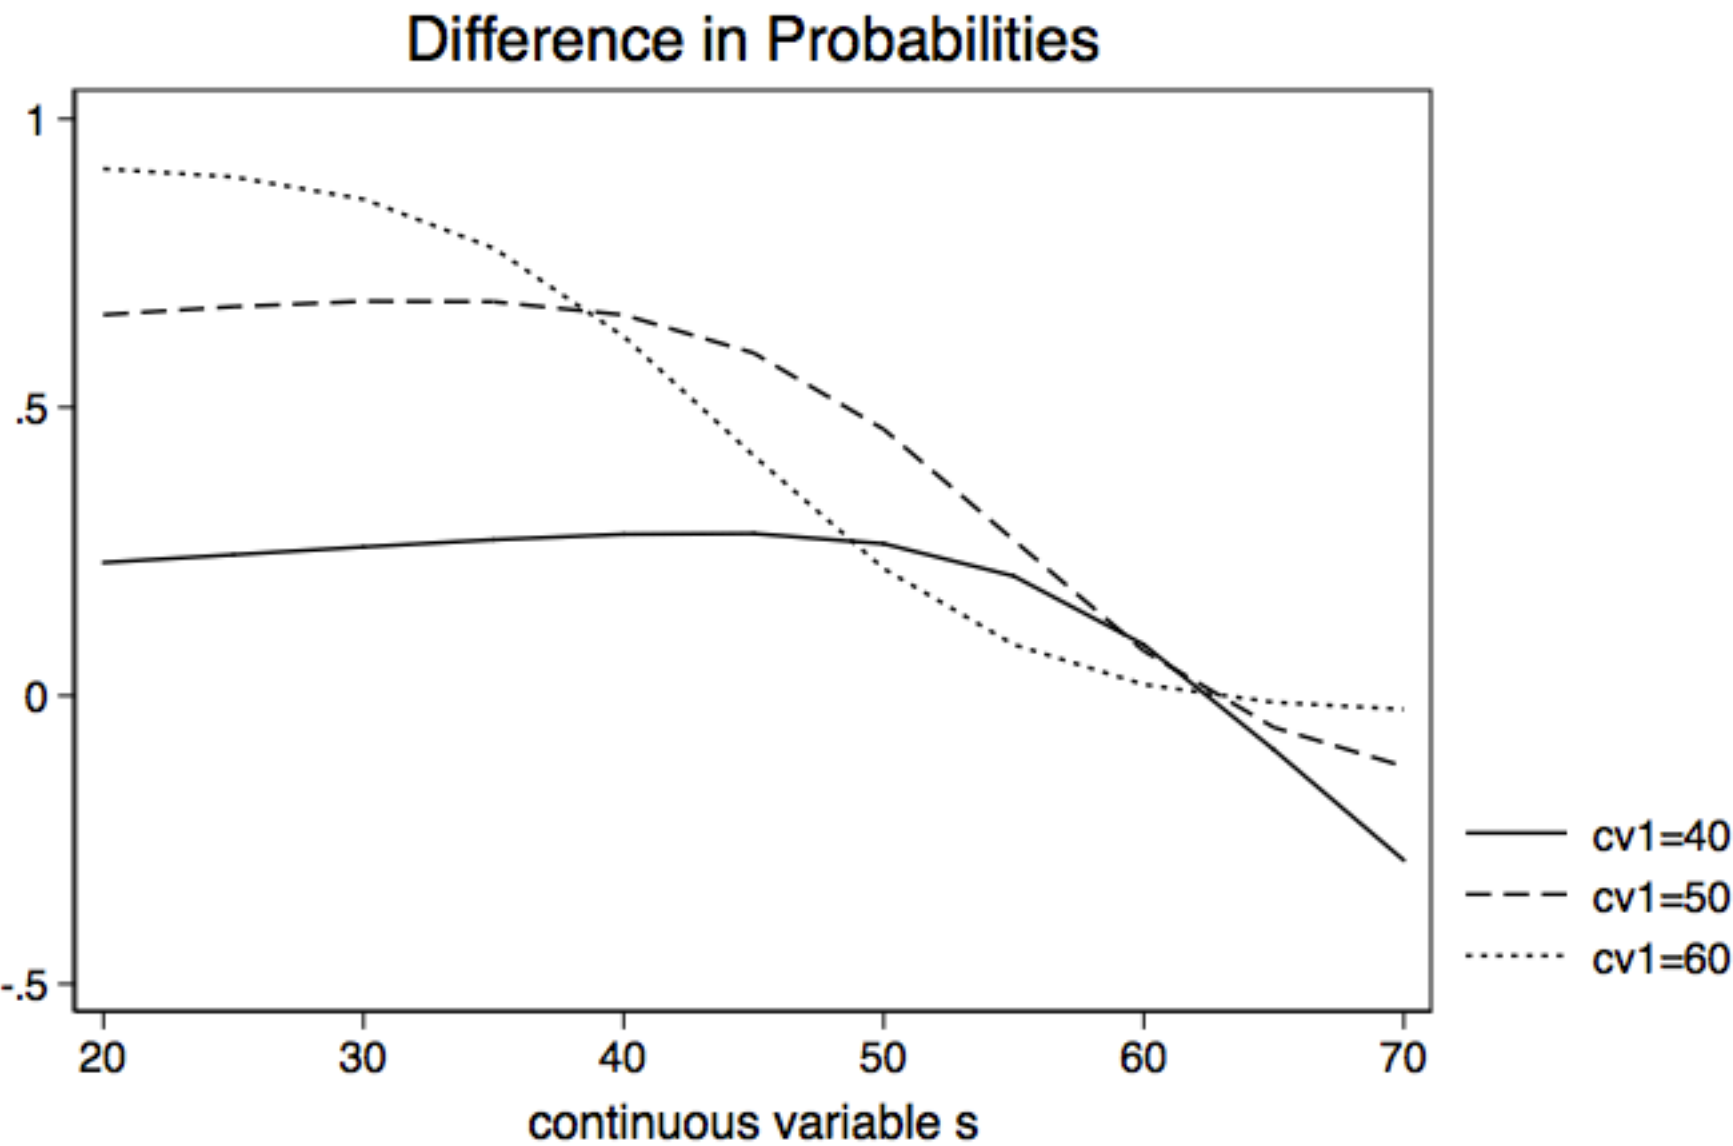
Instead of looking at separate values for **f0** and **f1**, we could compute the difference in probabilities. Here is an example using **margins** with the **dydx** option.

<b>margins, dydx(f) at(s=20 cv1=40)</b>						
Conditional marginal effects				Number of obs	=	200
Model VCE : OIM						
Expression : Pr(y), predict()						
dy/dx w.r.t. : 1.f						
at	:	s	=	20		
		cv1	=	40		
<hr/>						
		Delta-method				
		dy/dx	Std. Err.	z	P> z	[95% Conf. Interval]
<hr/>						
	1.f	.2307214	.150045	1.54	0.124	-.0633615 .5248042
<hr/>						
Note: dy/dx for factor levels is the discrete change from the base level.						

Okay, let’s repeat this for different values of **s**, producing the table below.


Table of Differences in Probability for Various Values of s Holding cv1 at 40						
s	Delta-method		z	P> z	[95% Conf. Interval]	
	dy/dx	Std. Err.				
20	.2307214	.150045	1.54	0.124	-.0633615	.5248042
25	.2443475	.1321009	1.85	0.064	-.0145655	.5032605
30	.2578855	.1135271	2.27	0.023	.0353765	.4803946
35	.2704118	.0954463	2.83	0.005	.0833405	.4574832
40	.2797622	.0798258	3.50	0.000	.1233066	.4362179
45	.2809858	.0696338	4.04	0.000	.1445061	.4174655
50	.2634265	.0682395	3.86	0.000	.1296795	.3971735
55	.2073734	.0822883	2.52	0.012	.0460913	.3686556
60	.08845	.1291224	0.69	0.493	-.1646253	.3415252
65	-.0937888	.2006804	-0.47	0.640	-.4871151	.2995376
70	-.2855934	.2436296	-1.17	0.241	-.7630986	.1919118
Note: dy/dx for factor levels is the discrete change from the base level.						

Next, we need to repeat the process while holding **cv1** at 50 and then 60. Then we can plot the differences in probabilities for the three values of **cv1** on a single graph.




The Stata FAQ page, [How can I understand a categorical by continuous interaction in logistic regression?](http://www.stata.com/faq/how-can-i-understand-a-categorical-by-continuous-interaction-in-logistic-regression-stata-11/) ([/stata/faq/how-can-i-understand-a-categorical-by-continuous-interaction-in-logistic-regression-stata-11/](http://www.stata.com/faq/how-can-i-understand-a-categorical-by-continuous-interaction-in-logistic-regression-stata-11/)) shows an alternative method for graphing these difference in probability lines to include confidence intervals. Here are the graphs from that FAQ page.

male-female difference with cv1 at 40



male-female difference with cv1 at 50





The trick to interpreting continuous by continuous interactions is to fix one predictor at a given value and to vary the other predictor. Once again, since the log odds model is a linear model it really doesn't matter what value the covariate is held at; the slopes do not change. For convenience we will just hold **cv1** at zero.

Here is an example manual computation of the slope of **r** holding **m** at 30.

<b>slope = b[r] + 30*b[r#m] = .43420626 + 30*(-.00681441) = .22977396</b>	
---	--

Here is the same computation using Stata.

<b>margins, dydx(r) at(m=30) predict(xb)</b>	
Average marginal effects	Number of obs = 200
Model VCE : OIM	
Expression : Linear prediction, predict(xb)	
dy/dx w.r.t. : r	
at : m = 30	
cv1 = 0	
-----	
	Delta-method
	dy/dx Std. Err. z P> z  [95% Conf. Interval]
-----+-----	
r	.2297741 .0982943 2.34 0.019 .0371207 .4224274
-----	

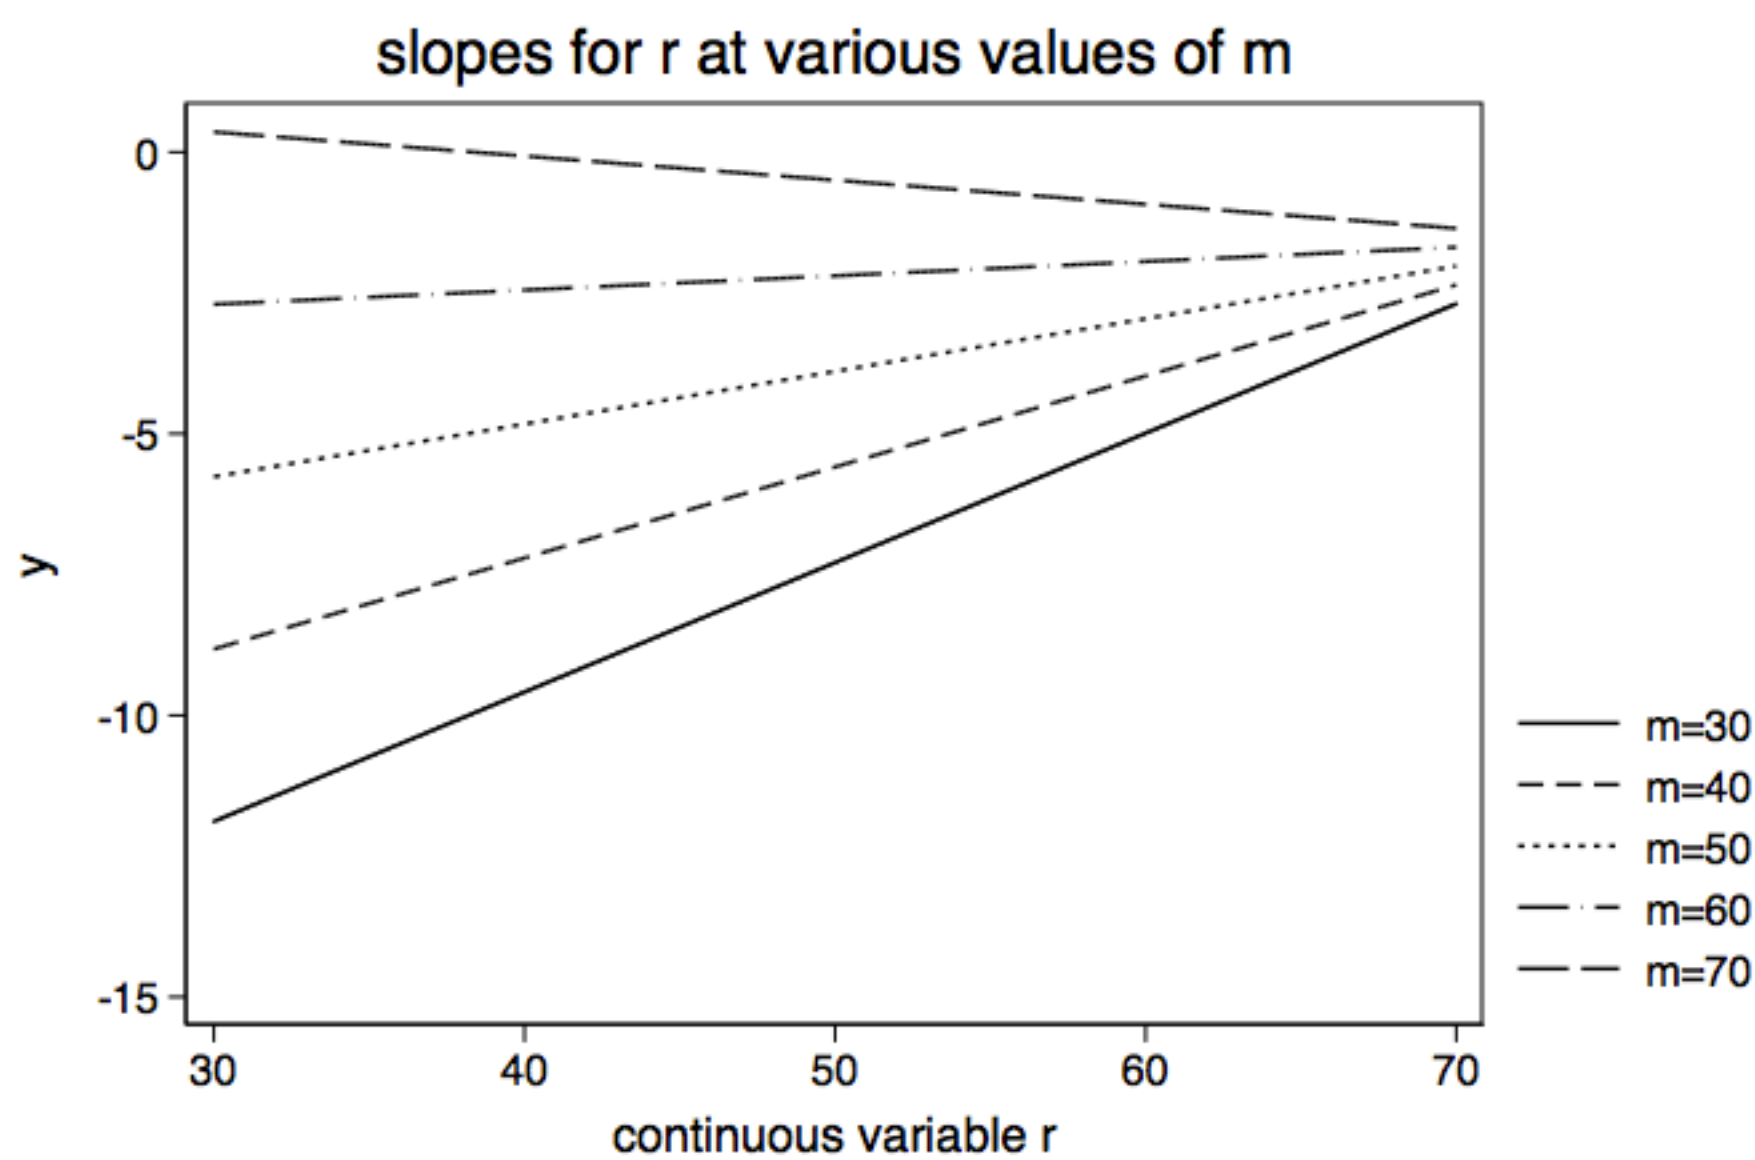
The table below shows the slope for **r** for various values of **m** running from 30 to 70. Since this is a linear model we do not have to hold **cv1** at any particular value.

<b>Table of Slopes for r for Various Values of m</b>	
	Delta-method
	dy/dx Std. Err. z P> z  [95% Conf. Interval]
-----+-----	
m	
30	.2297741 .0982943 2.34 0.019 .0371207 .4224274
40	.16163 .0670895 2.41 0.016 .0301369 .2931231
50	.0934859 .0395342 2.36 0.018 .0160004 .1709715
60	.0253419 .0291137 0.87 0.384 -.0317199 .0824037
70	-.0428022 .0485281 -0.88 0.378 -.1379156 .0523112
-----	



We arbitrarily chose to vary **m** and look at the slope of **r** but we could have easily reversed the variables. Hopefully, your knowledge of the theory behind the model along with substantive knowledge will suggest which variable to manipulate.

Below is a graph of the slopes from the table above.



This time we are going to move directly to the probability interpretation by-passing the odds ratio metric.

## Probability metric — continuous by continuous interaction

We will rerun our model.

```
logit y c.r##c.m cv1, nolog
```

Logistic regression	Number of obs	=	200
	LR chi2(4)	=	66.80
	Prob > chi2	=	0.0000
	Pseudo R2	=	0.3000
Log likelihood = -77.953857			

y	Coef.	Std. Err.	z	P> z	[95% Conf. Interval]	
r	.4342063	.1961642	2.21	0.027	.0497316	.8186809
m	.5104617	.2011856	2.54	0.011	.1161452	.9047782
c.r#c.m	-.0068144	.0033337	-2.04	0.041	-.0133483	-.0002805
cv1	.0309685	.0271748	1.14	0.254	-.0222931	.08423
_cons	-34.09122	11.73402	-2.91	0.004	-57.08947	-11.09297

Next we will calculate the values of the covariate for the mean minus one standard deviation, the mean, and the mean plus one standard deviation.

summarize cv1						
Variable		Obs	Mean	Std. Dev.	Min	Max
-----+-----						
cv1		200	52.405	10.73579	26	71
mean cv1 - 1sd = 41.669207						
mean cv1 = 52.405						
mean cv1 + 1sd = 63.140793						

Here is an example of a computation for the slope of r in the probability metric for **m** = 30 hold **cv1** at its mean minus 1 sd (41.669207).


margins, dydx(r) at(m=30 cv1=41.669207)						
Average marginal effects						
				Number of obs	=	200
Model VCE : OIM						
Expression : Pr(y), predict()						
dy/dx w.r.t. : r						
at : m = 30						
cv1 = 41.66921						
-----						
		Delta-method				
		dy/dx	Std. Err.	z	P> z	[95% Conf. Interval]
-----+-----						
r		.0061133	.0065712	0.93	0.352	-.006766 .0189926
-----						

We will now compute the slopes for r for differing values of **m** for each of the three values of **cv1**.


Table for Slope of r for Various Values of m holding cv1 at mean minus 1 sd							
	Delta-method						
	dy/dx	Std. Err.	z	P> z	[95% Conf. Interval]		
m							
30	.0061133	.0065712	0.93	0.352	-.006766	.0189926	
35	.006587	.0061377	1.07	0.283	-.0054427	.0186167	
40	.0071815	.0056839	1.26	0.206	-.0039586	.0183217	
45	.0078851	.0052656	1.50	0.134	-.0024354	.0182055	
50	.0085235	.004981	1.71	0.087	-.0012391	.0182861	
55	.0083341	.0049614	1.68	0.093	-.0013901	.0180583	
60	.0052692	.0059747	0.88	0.378	-.0064411	.0169795	
65	-.002175	.0090427	-0.24	0.810	-.0198984	.0155484	
70	-.0091967	.0089699	-1.03	0.305	-.0267774	.0083839	
Table for Slope of r for Various Values of m holding cv1 at the mean							
30	.0074917	.0069416	1.08	0.280	-.0061135	.0210969	
35	.0081075	.0063953	1.27	0.205	-.004427	.0206421	
40	.0088605	.0057648	1.54	0.124	-.0024384	.0201593	
45	.009721	.0051157	1.90	0.057	-.0003056	.0197476	
50	.0104242	.0046175	2.26	0.024	.0013739	.0194744	
55	.00992	.0046688	2.12	0.034	.0007692	.0190708	
60	.0058498	.006339	0.92	0.356	-.0065745	.0182741	
65	-.0021432	.0088189	-0.24	0.808	-.019428	.0151416	
70	-.0081533	.0075364	-1.08	0.279	-.0229243	.0066177	
Table for Slope of r for Various Values of m holding cv1 at mean plus 1 sd							
m							
30	.0090189	.0073769	1.22	0.221	-.0054396	.0234774	
35	.0097902	.0067546	1.45	0.147	-.0034485	.0230289	
40	.0107094	.0060155	1.78	0.075	-.0010807	.0224994	
45	.0117184	.0052384	2.24	0.025	.0014513	.0219854	
50	.0124196	.0046088	2.69	0.007	.0033864	.0214527	
55	.0114027	.004686	2.43	0.015	.0022182	.0205871	
60	.006181	.0067253	0.92	0.358	-.0070003	.0193622	
65	-.0020011	.0080879	-0.25	0.805	-.0178531	.0138509	
70	-.0069432	.0060361	-1.15	0.250	-.0187739	.0048874	

We will graph each of the three tables above.


marginal effect with cv1 held at mean minus 1sd



marginal effect with cv1 held at mean



marginal effect with cv1 held at mean plus 1sd



## The bottom line

- Just because the interaction term is significant in the log odds model, it doesn't mean that the probability difference in differences will be significant for values of the covariate of interest.
- Paradoxically, even if the interaction term is not significant in the log odds model, the probability difference in differences may be significant for some values of the covariate.
- In the probability metric the values of all the variables in the model matter.

## References


- Ai, C.R. and Norton E.C. 2003. Interaction terms in logit and probit models. *Economics Letters* 80(1): 123-129.
- Greenland, S. and Rothman, K.J. 1998. *Modern Epidemiology, 2nd Ed.* Philadelphia: Lippincott Williams and Wilkins.
- Mitchell, M.N. and Chen X. 2005. Visualizing main effects and interactions for binary logit model. *Stata Journal* 5(1): 64-82.
- Norton, E.C., Wang, H., and Ai, C. 2004 Computing interaction effects and standard errors in logit and probit models. *Stata Journal* 4(2): 154-167.

## Comma separated data files

**Categorical by categorical:** <https://stats.idre.ucla.edu/wp-content/uploads/2016/02/concon2.csv>  
(<https://stats.idre.ucla.edu/wp-content/uploads/2016/02/concon2.csv>)

**Categorical by continuous:** <https://stats.idre.ucla.edu/wp-content/uploads/2016/02/logitcatcon.csv>  
(<https://stats.idre.ucla.edu/wp-content/uploads/2016/02/logitcatcon.csv>)

**Continuous by continuous:** <https://stats.idre.ucla.edu/wp-content/uploads/2016/02/logitconcon.csv>  
(<https://stats.idre.ucla.edu/wp-content/uploads/2016/02/logitconcon.csv>)



[Click here to report an error on this page or leave a comment](#)

[How to cite this page \(https://stats.idre.ucla.edu/other/mult-pkg/faq/general/faq-how-do-i-cite-web-pages-and-programs-from-the-ucla-statistical-consulting-group/\)](https://stats.idre.ucla.edu/other/mult-pkg/faq/general/faq-how-do-i-cite-web-pages-and-programs-from-the-ucla-statistical-consulting-group/)