

# R Objects & Programmatic Data Manipulation

## Fundamental Techniques in Data Science



**Utrecht  
University**

Kyle M. Lang

Department of Methodology & Statistics  
Utrecht University

# Outline

---

## R Objects & Data Types

- Vectors & Matrices

- Lists & Data Frames

- Factors

## Programmatic Data Manipulation

- Subsetting

- Transforming

- Rearranging

- Pipes



# R OBJECTS & DATA TYPES



# Vectors

---

Vectors are the simplest kind of R object.

- There is no concept of a “scalar” in R.

Vectors come in one of six “atomic modes”:

- numeric/double
- logical
- character
- integer
- complex
- raw



# Vectors

---

```
(v1 <- vector("numeric", 3))
```

```
[1] 0 0 0
```

```
(v2 <- vector("logical", 3))
```

```
[1] FALSE FALSE FALSE
```

```
(v3 <- vector("character", 3))
```

```
[1] "" "" ""
```

```
(v4 <- vector("integer", 3))
```

```
[1] 0 0 0
```

```
(v5 <- vector("complex", 3))
```

```
[1] 0+0i 0+0i 0+0i
```

```
(v6 <- vector("raw", 3))
```

```
[1] 00 00 00
```

# Generating Vectors

---

We have many ways of generating vectors.

```
(y1 <- c(1, 2, 3))
```

```
[1] 1 2 3
```

```
(y2 <- c(TRUE, FALSE, TRUE, TRUE))
```

```
[1] TRUE FALSE TRUE TRUE
```

```
(y3 <- c("bob", "suzy", "danny"))
```

```
[1] "bob" "suzy" "danny"
```

```
1:5
```

```
[1] 1 2 3 4 5
```

```
1.2:5.3
```

```
[1] 1.2 2.2 3.2 4.2 5.2
```

# Generating Vectors

---

```
rep(33, 4)
```

```
[1] 33 33 33 33
```

```
rep(1:3, 3)
```

```
[1] 1 2 3 1 2 3 1 2 3
```

```
rep(y3, each = 2)
```

```
[1] "bob"    "bob"    "suzy"   "suzy"   "danny"  "danny"
```

```
seq(0, 1, 0.25)
```

```
[1] 0.00 0.25 0.50 0.75 1.00
```



# The Three Most Useful Data Types

---

## Numeric

```
(a <- 1:5)
[1] 1 2 3 4 5
```

## Character

```
(b <- c("foo", "bar"))
[1] "foo" "bar"
```

## Logical

```
(c <- c(TRUE, FALSE))
[1] TRUE FALSE
```





# Combining Data Types in Vectors

What happens if we try to concatenate different data types?

```
c(a, b)
```

```
[1] "1" "2" "3" "4" "5" "foo" "bar"
```

```
c(b, c)
```

```
[1] "foo" "bar" "TRUE" "FALSE"
```

```
c(a, c)
```

```
[1] 1 2 3 4 5 1 0
```



# Matrices

---

Matrices generalize vectors by adding a dimension attribute.

```
(m1 <- matrix(a, nrow = 5, ncol = 2))
```

	[,1]	[,2]
[1,]	1	1
[2,]	2	2
[3,]	3	3
[4,]	4	4
[5,]	5	5

```
attributes(v1)
```

```
NULL
```

```
attributes(m1)
```

```
$dim  
[1] 5 2
```

# Matrices

---

Matrices are populated in column-major order, by default.

```
(m2 <- matrix(1:9, 3, 3))
```

	[,1]	[,2]	[,3]
[1,]	1	4	7
[2,]	2	5	8
[3,]	3	6	9

The `byrow = TRUE` option allows us to fill by row-major order.

```
(m3 <- matrix(1:9, 3, 3, byrow = TRUE))
```

	[,1]	[,2]	[,3]
[1,]	1	2	3
[2,]	4	5	6
[3,]	7	8	9

# Mixing Data Types in Matrices

Like vectors, matrices can only hold one type of data.

```
cbind(c, letters[1:5])
```

```
      c
[1,] "TRUE"  "a"
[2,] "FALSE" "b"
[3,] "TRUE"  "c"
[4,] "FALSE" "d"
[5,] "TRUE"  "e"
```

```
cbind(c, c(TRUE, TRUE, FALSE, FALSE, TRUE))
```

```
      c
[1,] TRUE TRUE
[2,] FALSE TRUE
[3,] TRUE FALSE
[4,] FALSE FALSE
[5,] TRUE TRUE
```

# Lists

---

Lists are the workhorse of R data objects.

- An R list can hold an arbitrary set of other R objects.

We create lists using the `list()` function.

```
(l1 <- list(1, 2, 3))
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] 2
```

```
[[3]]
```

```
[1] 3
```

# Lists

---

```
(l2 <- list("bob", TRUE, 33, 42+3i))
```

```
[[1]]
```

```
[1] "bob"
```

```
[[2]]
```

```
[1] TRUE
```

```
[[3]]
```

```
[1] 33
```

```
[[4]]
```

```
[1] 42+3i
```



# Lists

---

List elements have no default names, but we can define our own.

```
(l3 <- list(name = "bob",  
            alive = TRUE,  
            age = 33,  
            relationshipStatus = 42+3i)  
)  
  
$name  
[1] "bob"  
  
$alive  
[1] TRUE  
  
$age  
[1] 33  
  
$relationshipStatus  
[1] 42+3i
```

# Lists

---

We can also assign post hoc names via the `names()` function.

```
names(l1) <- c("first", "second", "third")
l1

$first
[1] 1

$second
[1] 2

$third
[1] 3
```





# Lists

---

We can append new elements onto an existing list.

```
(l4 <- list())  
  
list()  
  
l4$people <- c("Bob", "Alice", "Suzy")  
l4$money <- 0  
l4$logical <- FALSE  
l4  
  
$people  
[1] "Bob" "Alice" "Suzy"  
  
$money  
[1] 0  
  
$logical  
[1] FALSE
```

# Lists

---

The elements inside a list don't really know that they live in a list; they'll pretty much behave as normal.

```
l4$money + 42
```

```
[1] 42
```

```
paste0("Hello, ", l4$people, "!\n") %>% cat()
```

```
Hello, Bob!
```

```
Hello, Alice!
```

```
Hello, Suzy!
```



# Data Frames

---

Data frames are R's way of storing rectangular data sets.

- Each column of a data frame is a vector.
- Each of these vectors can have a different type.

We create data frames using the `data.frame()` function.

```
(d1 <- data.frame(1:6, c(-1, 1), seq(0.1, 0.6, 0.1)))
```

	X1.6	c..1..1.	seq.0.1..0.6..0.1.
1	1	-1	0.1
2	2	1	0.2
3	3	-1	0.3
4	4	1	0.4
5	5	-1	0.5
6	6	1	0.6

# Data Frames

---

```
(d2 <- data.frame(x = 1:6, y = c(-1, 1), z = seq(0.1, 0.6, 0.1)))
```

	x	y	z
1	1	-1	0.1
2	2	1	0.2
3	3	-1	0.3
4	4	1	0.4
5	5	-1	0.5
6	6	1	0.6



# Data Frames

---

```
(d3 <- data.frame(a = sample(c(TRUE, FALSE), 10, replace = TRUE),  
                  b = sample(c("foo", "bar"), 10, replace = TRUE),  
                  c = runif(10)  
                  )
```

```
)
```

	a	b	c
1	FALSE	foo	0.4708232
2	TRUE	foo	0.2701596
3	TRUE	bar	0.6199154
4	FALSE	bar	0.2078104
5	TRUE	bar	0.4912943
6	FALSE	bar	0.1840306
7	FALSE	bar	0.5438698
8	TRUE	bar	0.5755350
9	FALSE	bar	0.7557042
10	FALSE	bar	0.8405729

# Data Frames

---

```
(d4 <- data.frame(matrix(NA, 10, 3)))
```

	X1	X2	X3
1	NA	NA	NA
2	NA	NA	NA
3	NA	NA	NA
4	NA	NA	NA
5	NA	NA	NA
6	NA	NA	NA
7	NA	NA	NA
8	NA	NA	NA
9	NA	NA	NA
10	NA	NA	NA



# Data Frames

---

Data frames are actually lists of vectors (representing the columns).

```
is.data.frame(d3)
```

```
[1] TRUE
```

```
is.list(d3)
```

```
[1] TRUE
```

Although they look like rectangular “matrices”, from R’s perspective a data frame IS NOT a matrix.

```
is.matrix(d3)
```

```
[1] FALSE
```

# Data Frames

---

We cannot treat a data frame like a matrix. E.g., matrix algebra doesn't work with data frames.

```
d1 %*% t(d2)
```

```
Error in d1 %*% t(d2): requires numeric/complex matrix/vector arguments
```

```
as.matrix(d1) %*% t(as.matrix(d2))
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]
[1,]	2.01	1.02	4.03	3.04	6.05	5.06
[2,]	1.02	5.04	5.06	9.08	9.10	13.12
[3,]	4.03	5.06	10.09	11.12	16.15	17.18
[4,]	3.04	9.08	11.12	17.16	19.20	25.24
[5,]	6.05	9.10	16.15	19.20	26.25	29.30
[6,]	5.06	13.12	17.18	25.24	29.30	37.36



# Factors

---

Factors are R's way of representing nominal variables.

- We can create a factor using the `factor()` function.

```
(f1 <- factor(sample(1:3, 15, TRUE), labels = c("red", "yellow", "blue")))  
[1] red    yellow blue   red    blue   blue   yellow red  
[9] red    red    yellow blue   blue   blue   yellow  
Levels: red yellow blue
```



# Factors

---

Factors are integer vectors with a *levels* attribute and a *factor* class.

```
typeof(f1)

[1] "integer"

attributes(f1)

$levels
[1] "red"      "yellow" "blue"

$class
[1] "factor"
```

The levels are just group labels.

```
levels(f1)

[1] "red"      "yellow" "blue"
```

# Factors

---

Even though a factor's data are represented by an integer vector, R does not consider factors to be integer/numeric data.

```
is.numeric(f1)
```

```
[1] FALSE
```

```
is.integer(f1)
```

```
[1] FALSE
```

Factors represent nominal variables, so we cannot do math with factors.

```
f1 + 1
```

```
[1] NA NA NA NA NA NA NA NA NA NA NA NA NA NA NA
```

```
mean(f1)
```

```
[1] NA
```

# PROGRAMMATIC DATA MANIPULATION



# Base R Subsetting

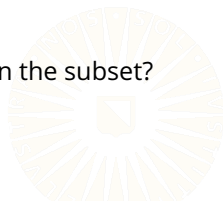
---

In Base R, we typically use three operators to subset objects:

- `[]`
- `[[ ]]`
- `$`

Which of these operators we choose to use (and how we implement the chosen operator) will depend on two criteria:

- What type of object are we trying to subset?
- How much of the original typing do we want to keep in the subset?



# Tidyverse Subsetting

---

The **dplyr** package provides many ways to subset data, but two functions are most frequently useful.

- `select()` : subset columns
- `filter()` : subset rows

```
library(dplyr)
```



# Tidyverse Solutions

---

The **dplyr** package provides three primary transformation functions

- `recode()` : recode the levels of a variable
- `mutate()` : general purpose transformation & feature building

```
library(dplyr)
```



# Tidyverse Solutions

---

The **dplyr** package provides three primary transformation functions

- `arrange()` : sort/order rows
- `relocate()` : move columns

```
library(dplyr)
```





# What are pipes?

---

The %>% symbol represents the *pipe* operator.

- We use the pipe operator to compose functions into a *pipeline*.

The following code represents a pipeline.

```
firstBoys <-  
  read_sav("../data/boys.sav") %>%  
  head()
```

This pipeline replaces the following code.

```
firstBoys <- head(read_sav("../data/boys.sav"))
```



# Why are pipes useful?

---

Let's assume that we want to:

1. Load data
2. Transform a variable
3. Filter cases
4. Select columns

Without a pipe, we may do something like this:

```
library(haven)
library(dplyr)

boys <- read_sav("../data/boys.sav")
boys <- transform(boys, hgt = hgt / 100)
boys <- filter(boys, age > 15)
boys <- subset(boys, select = c(hgt, wgt, bmi))
```

# Why are pipes useful?

---

With the pipe, we could do something like this:

```
library(haven)
library(dplyr)

boys <-
  read_sav("../..data/boys.sav") %>%
  transform(hgt = hgt / 100) %>%
  filter(age > 15) %>%
  subset(select = c(hgt, wgt, bmi))
```

With a pipeline, our code more clearly represents the sequence of steps in our analysis.



# Benefits of Pipes

---

When you use pipes, your code becomes more readable.

- Operations are structured from left to right instead of in to out.
- You can avoid many nested function calls.
- You don't have to keep track of intermediate objects.
- It's easy to add steps to the sequence.

In RStudio, you can use a keyboard shortcut to insert the `%>%` symbol.

- Windows/Linux: *ctrl + shift + m*
- Mac: *cmd + shift + m*



# What do pipes do?

---

Pipes compose R functions without nesting.

- `f(x)` becomes `x %>% f()`

```
mean(rnorm(10))
```

```
[1] 0.8904119
```

```
rnorm(10) %>% mean()
```

```
[1] 0.293541
```



# What do pipes do?

Multiple function arguments are fine.

- `f(x, y)` becomes `x %>% f(y)`

```
cor(boys, use = "pairwise.complete.obs")
```

	hgt	wgt	bmi
hgt	1.0000000	0.6100784	0.1758781
wgt	0.6100784	1.0000000	0.8841304
bmi	0.1758781	0.8841304	1.0000000

```
boys %>% cor(use = "pairwise.complete.obs")
```

	hgt	wgt	bmi
hgt	1.0000000	0.6100784	0.1758781
wgt	0.6100784	1.0000000	0.8841304
bmi	0.1758781	0.8841304	1.0000000

# What do pipes do?

---

Composing more than two functions is easy, too.

- `h(g(f(x)))` becomes `x %>% f %>% g %>% h`

```
max(na.omit(subset(boys, select = wgt)))
```

```
[1] 117.4
```

```
boys %>%
```

```
  subset(select = wgt) %>%
```

```
  na.omit() %>%
```

```
  max()
```

```
[1] 117.4
```



# The Role of `.` in a Pipeline

In the expression `a %>% f(arg1, arg2, arg3)`, `a` will be "piped into" `f()` as `arg1`.

```
data(cats, package = "MASS")  
cats %>% plot(Hwt ~ Bwt)
```

```
Error in text.default(x, y, txt, cex = cex, font = font): invalid  
mathematical annotation
```

Clearly, we have a problem if we pipe our data into the wrong argument.


- We can change this behavior with the `.` symbol.
- The `.` symbol acts as a placeholder for the data in a pipeline.





# The Role of `.` in a Pipeline

```
cats %>% plot(Hwt ~ Bwt, data = .)
```




# Exposition Pipe: %\$%

There are several different flavors of pipe. The *exposition pipe*, %\$%, is a particularly useful variant.

- The exposition pipe *exposes* the contents of an object to the next function in the pipeline.

```
cats %$% plot(Hwt ~ Bwt)
```



# Performing a T-Test in a Pipeline

```
cats %$% t.test(Hwt ~ Sex)
```

Welch Two Sample t-test

data: Hwt by Sex

t = -6.5179, df = 140.61, p-value = 1.186e-09

alternative hypothesis: true difference in means is not equal to 0

95 percent confidence interval:

-2.763753 -1.477352

sample estimates:

mean in group F mean in group M

9.202128

11.322680

The above is equivalent to either of the following.

```
cats %>% t.test(Hwt ~ Sex, data = .)
```

```
t.test(Hwt ~ Sex, data = cats)
```

# Storing the Results

---

We can use normal assignment to save the result of a pipeline.

```
catsTest <- cats %>% t.test(Bwt ~ Sex)

catsTest

Welch Two Sample t-test

data:  Bwt by Sex
t = -8.7095, df = 136.84, p-value = 8.831e-15
alternative hypothesis: true difference in means is not equal to 0
95 percent confidence interval:
 -0.6631268 -0.4177242
sample estimates:
mean in group F mean in group M
      2.359574      2.900000
```

# References

---

