

# Data Cleaning, Data Visualization, & Functions

## Fundamental Techniques in Data Science



**Utrecht  
University**

Kyle M. Lang

Department of Methodology & Statistics  
Utrecht University

# Outline

---

## Functions

## Data Visualization

- Base R Graphics
- GGPlot

## Data Cleaning

- Data Analytic Lifecycle
- Missing Data
- Outliers



# FUNCTIONS



# R Functions

---

Functions are the foundation of R programming.

- Other than data objects, almost everything else that you interact with when using R is a function.
- Any R command written as a word followed by parentheses `()` is a function.
  - `mean()`
  - `library()`
  - `mutate()`
- Infix operators are aliased functions.
  - `<-`
  - `+`, `-`, `*`
  - `%>%`, `$$`, `%<>%`



# User-Defined Functions

---

We can define our own functions using the `function()` function.

```
square <- function(x) {  
  out <- x^2  
  out  
}
```

After defining a function, we use it in the same way as any other R function.

```
square(5)  
[1] 25
```



# User-Defined Functions

One-line functions don't need braces

```
square <- function(x) x^2
```

```
square(5)
```

```
[1] 25
```

Function arguments are not strictly typed. R will try to work with whatever you provide as input.

```
square(1:5)
```

```
[1] 1 4 9 16 25
```

```
square(pi)
```

```
[1] 9.869604
```

```
square(TRUE)
```

```
[1] 1
```

```
square("bob") # But one can only try so hard
```

```
Error in x^2
```

# User-Defined Functions

Functions can take multiple arguments.

```
mod <- function(x, y) x %% y
mod(10, 3)

[1] 1
```

Sometimes it's useful to specify a list of arguments that we unpack inside the function.

```
getLsBeta <- function(datList) {
  X <- datList$X
  y <- datList$y

  solve(crossprod(X)) %*% t(X) %*% y
}

X <- matrix(runif(500), ncol = 5)
datList <- list(y = X %*% rep(0.5, 5), X = X)

getLsBeta(datList = datList)

[,1]
[1,] 107 0.5
```

# User-Defined Functions

---

Functions are first-class objects in R.

- We can treat them like any other R object.

R views an initialized, but unevaluated, function as a special object with type "closure"

```
class(getLsBeta)
[1] "function"

typeof(getLsBeta)
[1] "closure"
```

After evaluation, functions are simply equivalent to the objects they return.

```
class(getLsBeta(datList))
[1] "matrix" "array"

typeof(getLsBeta(datList))
[1] "double"
```



# User-Defined Functions

We can use functions as arguments to other operations and functions.

```
fun1 <- function(x, y) x + y

## What will this command return?
fun1(1, fun1(1, 1))

[1] 3
```

Why would we care?

```
s2 <- var(runif(100))
x <- rnorm(100, 0, sqrt(s2))

x[1:10, ]

      [,1]      [,2]      [,3]      [,4]      [,5]
[1,] 0.52431382 0.67136447 0.28228726 0.7148383 0.54204681
[2,] 0.01926742 0.11693762 0.09148502 0.6929171 0.88371944
[3,] 0.05100735 0.18432074 0.43547799 0.6097462 0.09026598
[4,] 0.60566972 0.12944127 0.21000143 0.2441917 0.68141473
[5,] 0.48737303 0.94030405 0.23988619 0.4915910 0.36353771
[6,] 0.19941958 0.96670678 0.11455820 0.1243947 0.24253273
[7,] 0.95507004 0.88705000 0.48788585 0.8888170 0.81881888
[8,] 0.66107 0.95507004 0.88705000 0.48788585 0.8888170 0.81881888
```

# DATA VISUALIZATION



# Setup

---

```
dataDir <- "../data/"  
figDir  <- "../figures/"
```

```
diabetes <- readRDS(paste0(dataDir, "diabetes.rds"))  
titanic  <- readRDS(paste0(dataDir, "titanic.rds"))
```

```
Error in gzfile(file, "rb"): cannot open the connection
```

```
bfi      <- readRDS(paste0(dataDir, "bfi.rds"))
```

```
Error in gzfile(file, "rb"): cannot open the connection
```

```
## Convert survival indicator to a numeric dummy code:
```

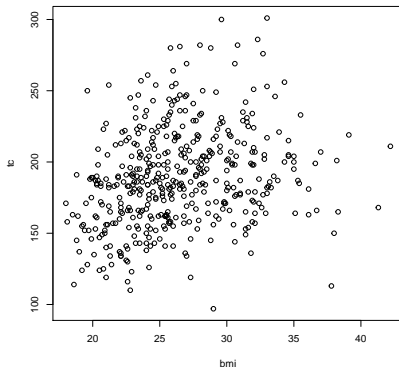
```
titanic <- titanic %>% mutate(survived = as.numeric(survived) - 1)
```

```
Error in mutate(., survived = as.numeric(survived) - 1): object 'titanic'  
not found
```

# Base R Graphics: Scatterplots

We can create a basic scatterplot using the `plot()` function.

```
diabetes %>% plot(y = tc, x = bmi)
```



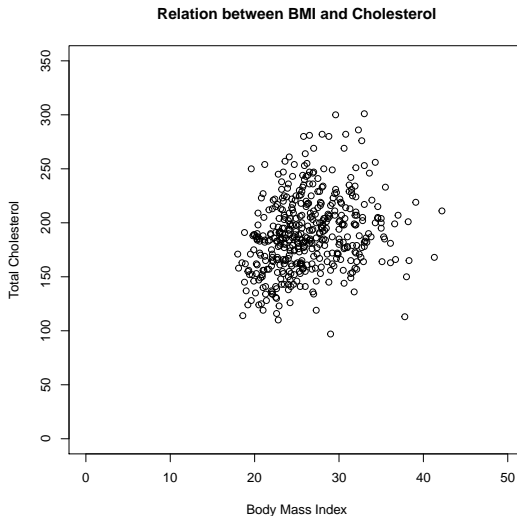
# Base R Graphics: Scatterplots

---

```
diabetes %$% plot(y = tc,  
                 x = bmi,  
                 ylab = "Total Cholesterol",  
                 xlab = "Body Mass Index",  
                 main = "Relation between BMI and Cholesterol",  
                 ylim = c(0, 350),  
                 xlim = c(0, 50)  
                 )
```



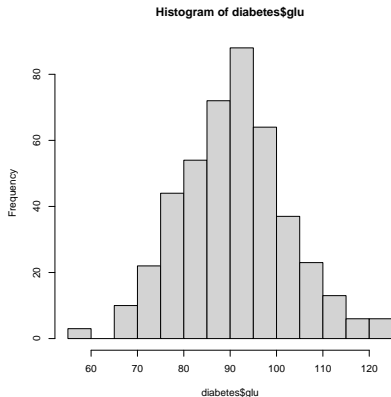
# Base R Graphics: Scatterplots



# Base R Graphics: Histograms

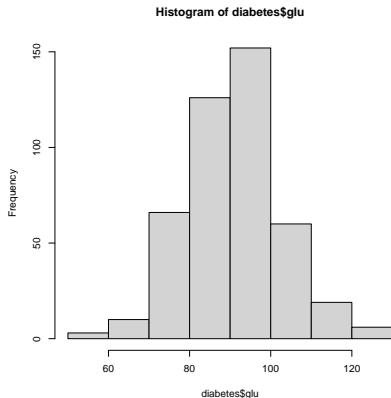
We can create a simple histogram with the `hist()` function.

```
hist(diabetes$glu)
```



# Base R Graphics: Histograms

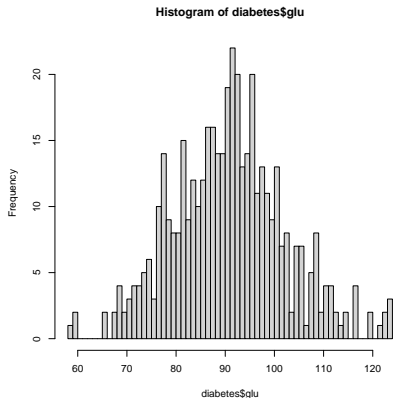
```
hist(diabetes$glu, breaks = 5)
```





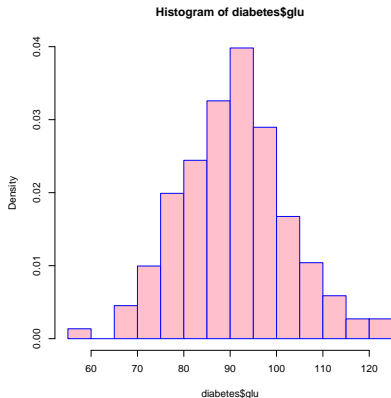
# Base R Graphics: Histograms

```
hist(diabetes$glu, breaks = 50)
```



# Base R Graphics: Histograms

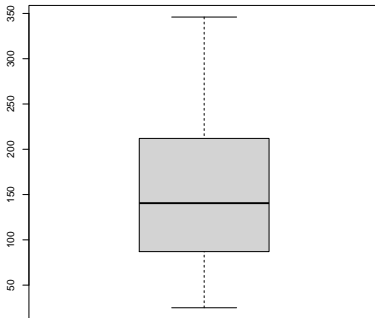
```
hist(diabetes$glu, col = "pink", border = "blue", probability = TRUE)
```



# Base R Graphics: Boxplots

We can create simple boxplots via the `boxplot()` function.

```
boxplot(diabetes$progress)
```



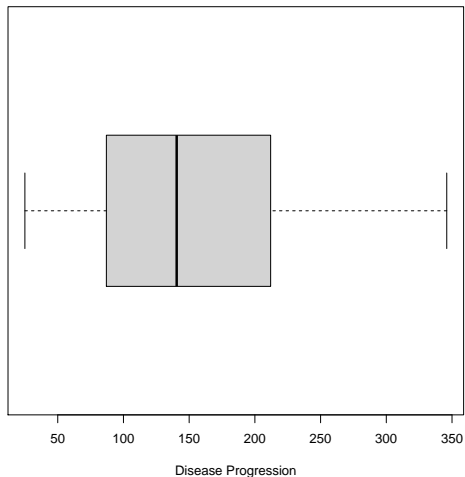
# Base R Graphics: Boxplots

---

```
boxplot(diabetes$progress,  
        horizontal = TRUE,  
        range = 3,  
        xlab = "Disease Progression")
```



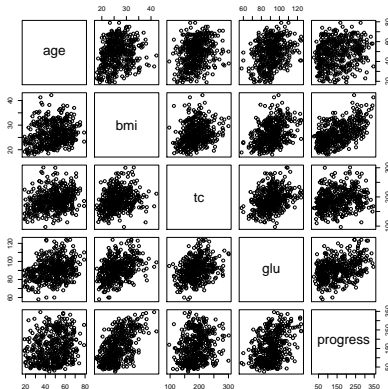
# Base R Graphics: Boxplots



# Base R Graphics: Fancy Things

Plotting an entire data frame produces a scatterplot matrix.

```
diabetes %>% select(age, bmi, tc, glu, progress) %>% plot()
```

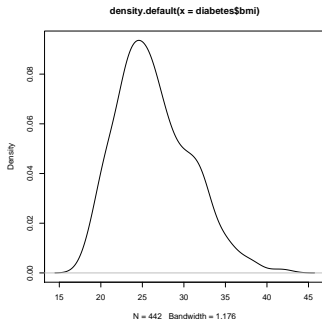


# Base R Graphics: Fancy Things

The `density()` function estimates the density of a variable.

- If we plot a density object, we get a kernel density plot.

```
density(diabetes$bmi) %>% plot()
```



# Base R Graphics: Fancy Things

---

```
d <- density(diabetes$bmi)
```

```
ls(d)
```

```
[1] "bw"          "call"        "data.name"  "has.na"
```

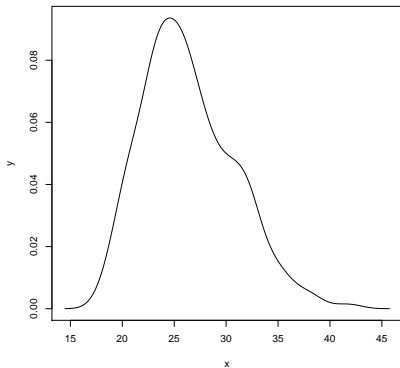
```
[5] "n"           "x"           "y"
```





# Base R Graphics: Fancy Things

```
d %>% plot(y = y, x = x, type = "l")
```



# Base R Graphics: Workflow

Base R graphics work by building up graphics from layers.

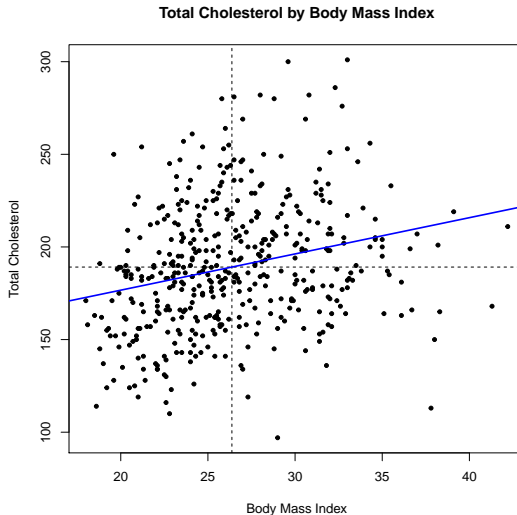
```
## Start with a simple scatterplot:
diabetes %$% plot(y = tc, x = bmi, pch = 20, xlab = "", ylab = "")

## Use the abline() function to add lines representing the means of x and y:
abline(h = mean(diabetes$tc), v = mean(diabetes$bmi), lty = 2)

## Add the best fit line from a linear regression of 'tc' onto 'bmi':
diabetes %$%
  lm(tc ~ bmi) %>%
  coef() %>%
  abline(coef = ., col = "blue", lwd = 2)

## Add titles:
title(main = "Total Cholesterol by Body Mass Index",
      ylab = "Total Cholesterol",
      xlab = "Body Mass Index")
```

# Base R Graphics: Workflow



# Base R Graphics: Workflow

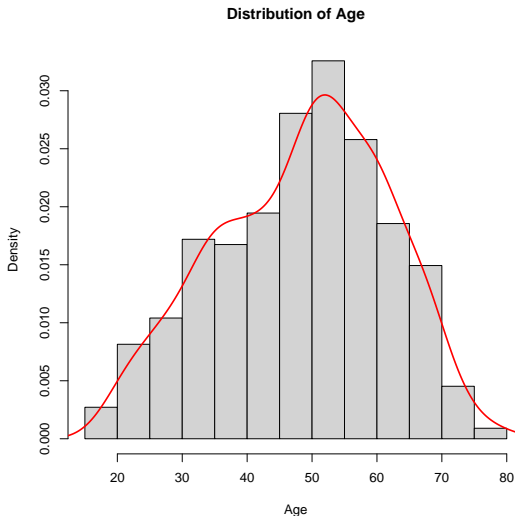
---

Add a kernel density plot on top of a histogram.

```
diabetes %$%  
  hist(age,  
        probability = TRUE,  
        xlab = "Age",  
        main = "Distribution of Age")  
  
diabetes %$%  
  density(age) %>%  
  lines(col = "red", lwd = 2)
```



# Base R Graphics: Workflow



# GGPlot

---

Base R graphics are fine for quick-and-dirty visualizations (e.g., EDA, checking assumptions), but for publication quality graphics, we probably want to use GGPlot.

GGPlot uses the "grammar of graphics" and "tidy data" to build up a figure from modular components

Describes all the non-data ink  
Plotting space for the data  
Statistical models & summaries  
Rows and columns of sub-plots  
Shapes used to represent the data  
Scales onto which data is mapped  
The actual variables to be plotted

**Theme**  
**Coordinates**  
**Statistics**  
**Facets**  
**Geometries**  
**Aesthetics**  
**Data**



# GGPlot: Basic Setup

---

We start by calling the `ggplot()` function.

- We must define a data source.
- We must also give some aesthetic via the `aes()` function.

```
library(ggplot2)
p1 <- ggplot(data = diabetes, mapping = aes(x = bmi, y = glu))
```

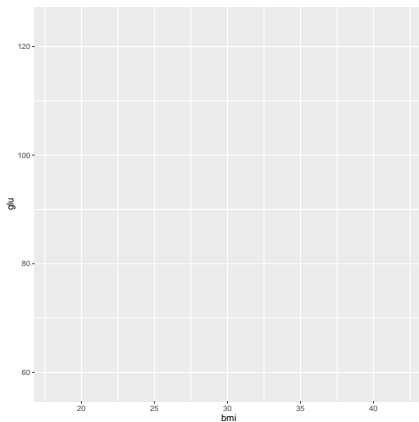


# GGPlot: Basic Setup

---

At this point, our plot is pretty boring.

p1





# GGPlot: Geometries

---

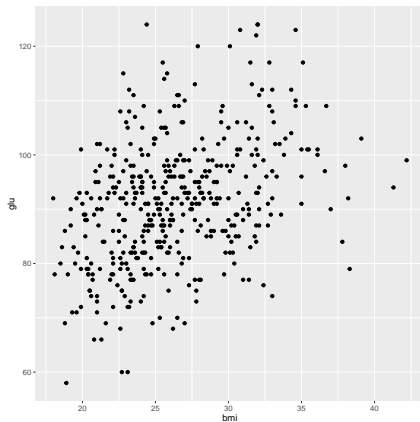
We need to define some geometry via an appropriate `geom_X()` function.

```
p1 + geom_point()
```




# GGPlot: Geometries

---



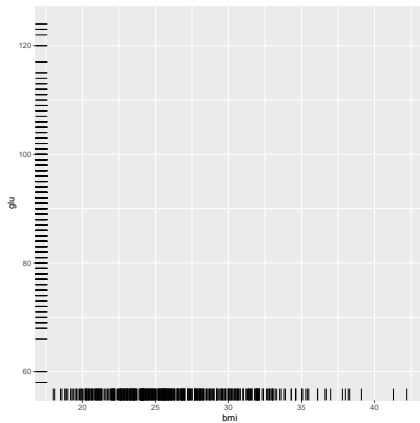
# GGPlot: Geometries

```
p1 + geom_line()
```



# GGPlot: Geometries

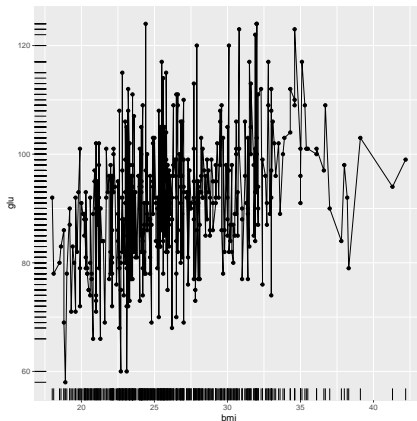
```
p1 + geom_rug()
```



# GGPlot: Geometries

We can also combine different geoms into a single figure

```
p1 + geom_point() + geom_line() + geom_rug()
```



# GGPlot: Geometries

---

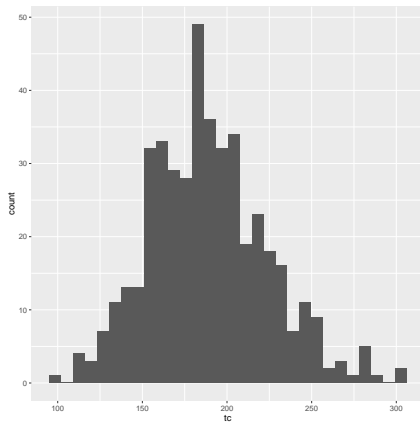
We can use different flavors of geom for different types of data

```
p2 <- ggplot(diabetes, aes(tc))  
p2 + geom_histogram()
```



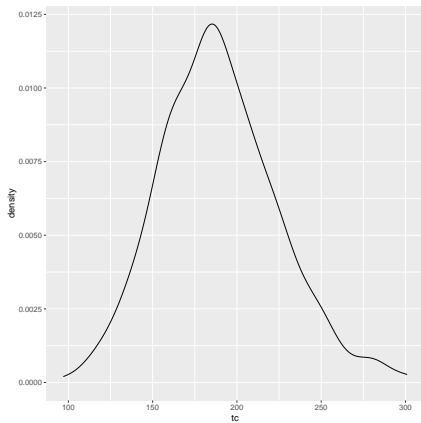
# GGPlot: Geometries

---



# GGPlot: Geometries

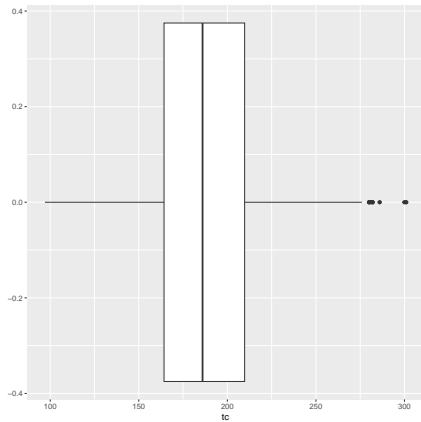
```
p2 + geom_density()
```





# GGPlot: Geometries

```
p2 + geom_boxplot()
```



# GGPlot: Geometries

---

```
p3 <- ggplot(diabetes, aes(sex, bmi))  
p3 + geom_boxplot()
```

```
Error in 'geom_boxplot()':  
! Problem while computing aesthetics.  
i Error occurred in the 1st layer.  
Caused by error in 'FUN()':  
! object 'sex' not found
```



# GGPlot: Geometries

```
p3 + geom_violin()
```

```
Error in 'geom_violin()':
! Problem while computing aesthetics.
i Error occurred in the 1st layer.
Caused by error in 'FUN()':
! object 'sex' not found
```

# GGPlot: Geometries

---

```
p4 <- ggplot(bfi, aes(education, age))
```

```
Error in ggplot(bfi, aes(education, age)): object 'bfi' not found
```

```
p4 + geom_point()
```

```
Error in eval(expr, envir, enclos): object 'p4' not found
```



# GGPlot: Geometries

---

```
p4 + geom_jitter()
```

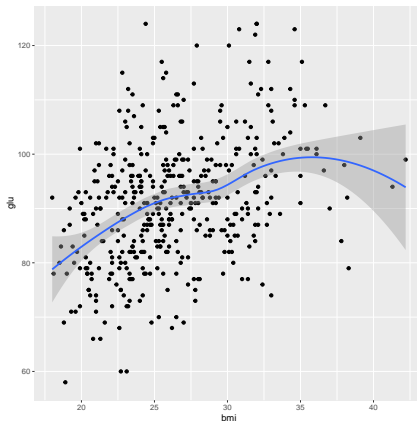
```
Error in eval(expr, envir, enclos): object 'p4' not found
```



# GGPlot: Statistics

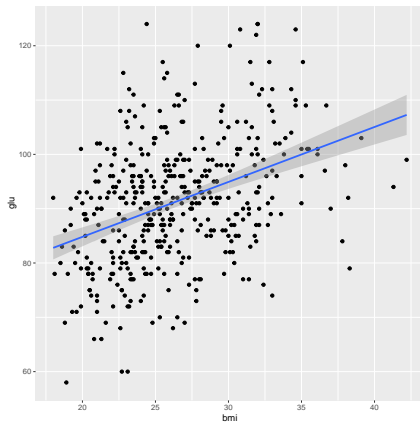
We can also add statistical summaries of the data

```
p1 + geom_point() + geom_smooth()
```



# GGPlot: Statistics

```
p1 + geom_point() + geom_smooth(method = "lm")
```



# GGPlot: Styling

---

Changing style options outside of the `aes()` function applies the styling to the entire plot.

```
p5 <- ggplot(titanic, aes(age, survived))
```

```
Error in ggplot(titanic, aes(age, survived)): object 'titanic' not found
```

```
p5 + geom_jitter(color = "blue", size = 3, height = 0.1)
```

```
Error in eval(expr, envir, enclos): object 'p5' not found
```





# GGPlot: Styling

---

We can also apply styles as a function of variables by defining the style within the `aes()` function.

```
p6.1 <- ggplot(titanic, aes(age, survived, color = sex))
```

```
Error in ggplot(titanic, aes(age, survived, color = sex)): object  
'titanic' not found
```

```
p6.1 + geom_jitter(size = 3, height = 0.1) + geom_smooth()
```

```
Error in eval(expr, envir, enclos): object 'p6.1' not found
```



# GGPlot: Styling

---

```
p6.2 <- ggplot(titanic, aes(age, survived))
```

```
Error in ggplot(titanic, aes(age, survived)): object 'titanic' not found
```

```
p6.2 + geom_jitter(aes(color = sex), size = 3, height = 0.1) +  
  geom_smooth()
```

```
Error in eval(expr, envir, enclos): object 'p6.2' not found
```



# GGPlot: Styling

---

```
p6.2 + geom_jitter(size = 3, height = 0.1) +  
  geom_smooth(aes(color = sex))
```

Error in eval(expr, envir, enclos): object 'p6.2' not found

```
p6.2 + geom_jitter(aes(color = class), size = 3, height = 0.1) +  
  geom_smooth(aes(color = sex))
```

Error in eval(expr, envir, enclos): object 'p6.2' not found



# GGPlot: Styling

---

```
p6.2 + geom_jitter(aes(shape = class), size = 3, height = 0.1) +  
  geom_smooth(aes(color = sex))
```

Error in eval(expr, envir, enclos): object 'p6.2' not found



# GGPlot: Styling

---

```
p6.1 + geom_jitter(aes(shape = class), size = 3, height = 0.1) +  
  geom_smooth()
```

Error in eval(expr, envir, enclos): object 'p6.1' not found



# GGPlot: Themes

---


We can apply several pre-baked themes to adjust a plot's overall appearance

```
(p1.1 <- p1 + geom_point())
```



# GGPlot: Themes


---



# GGPlot: Themes

---

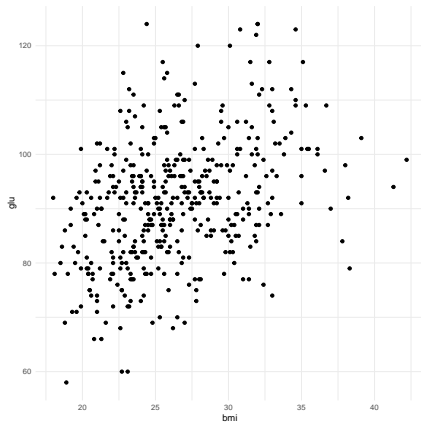
```
p1.1 + theme_classic()
```






# GGPlot: Themes

```
p1.1 + theme_minimal()
```



# GGPlot: Themes

```
p1.1 + theme_bw()
```



# GGPlot: Themes

---


We can also moodifying individual theme elements.

```
p1.1 + theme_classic() +  
  theme(axis.title = element_text(size = 16,  
                                   family = "serif",  
                                   face = "bold",  
                                   color = "blue"),  
        aspect.ratio = 1)
```



# GGPlot: Themes

---



# GGPlot: Facets

---

Facetting allow us to make arrays of conditional plots.

```
(p7 <- ggplot(titanic, aes(age, survived, color = class)) +  
  geom_jitter(height = 0.05) +  
  geom_smooth(method = "glm",  
              method.args = list(family = "binomial"),  
              se = FALSE)  
)
```



# GGPlot: Facets

---

```
Error in ggplot(titanic, aes(age, survived, color = class)): object  
'titanic' not found
```



# GGPlot: Facets

---

```
p7 + facet_wrap(vars(sex))
```

```
Error in eval(expr, envir, enclos): object 'p7' not found
```



# GGPlot: Facets

---

```
## Use facet_grid() to condition plots on both 'sex' and 'class'  
(p8 <- ggplot(titanic, aes(age, survived)) +  
  geom_jitter(height = 0.05) +  
  geom_smooth(method = "glm",  
              method.args = list(family = "binomial"),  
              se = FALSE)  
)
```





# GGPlot: Facets

---

```
Error in ggplot(titanic, aes(age, survived)): object 'titanic' not found
```



# GGPlot: Facets

---

```
Error in eval(expr, envir, enclos): object 'p8' not found
```



# GGPlot: Joining Multiple Figures

---

If we want to paste several different plots into a single figure (without facetting), we can use the utilities in the **gridExtra** package.

```
library(gridExtra)

grid.arrange(p1 + geom_point(),
             p3 + geom_boxplot(),
             p4 + geom_jitter(),
             p8 + facet_grid(vars(sex), vars(class)),
             ncol = 2)
```



# GGPlot: Joining Multiple Figures

---

```
Error in arrangeGrob(...): object 'p4' not found
```



# Saving Graphics

---

To save a graphic that we've created in R, we simply redirect the graphical output to a file using an appropriate function.

```
figDir <- "../figures/"

## Save as PDF
pdf(paste0(figDir, "example_plot.pdf"))

Error in pdf(paste0(figDir, "example_plot.pdf")): cannot open file
'../figures/example_plot.pdf'

p7 + facet_wrap(vars(sex))

Error in eval(expr, envir, enclos): object 'p7' not found

dev.off()

null device
      1
```

# Saving Graphics

---

```
## Save as JPEG
jpeg(paste0(figDir, "example_plot.jpg"))

p7 + facet_wrap(vars(sex))

Error in eval(expr, envir, enclos): object 'p7' not found

dev.off()

pdf
2
```



# Saving Graphics

---

```
## Save as PNG
png(paste0(figDir, "example_plot.png"))

p7 + facet_wrap(vars(sex))

Error in eval(expr, envir, enclos): object 'p7' not found

dev.off()

pdf
2
```



# Saving Graphics

---

With PDF documents, we can save multiple figures to a single file.





# Saving Graphics

---

```
pdf(paste0(figDir, "example_plot2.pdf"))
```

```
Error in pdf(paste0(figDir, "example_plot2.pdf")): cannot open file  
'../figures/example_plot2.pdf'
```

```
p6.1 + geom_jitter(size = 3, height = 0.1) + geom_smooth()
```

```
Error in eval(expr, envir, enclos): object 'p6.1' not found
```

```
p7 + facet_wrap(vars(sex))
```

```
Error in eval(expr, envir, enclos): object 'p7' not found
```

```
p8 + facet_grid(vars(sex), vars(class))
```

```
Error in eval(expr, envir, enclos): object 'p8' not found
```

```
dev.off()
```

```
null device  
1
```


# DATA CLEANING



# Research Cycle


---

The following is a representation of the *Research Cycle* used for empirical research in most of the sciences.




# CRISP-DM

The *Cross-industry Standard Process for Data Mining* was developed to standardized the process of data mining in industry applications.



# Data Science Cycle

The *Data Science Cycle* represented here was adapted from O'Neil and Schutt (2014).



# Data Cleaning

---

When we receive new data, they are generally messy and contaminated by various anomalies and errors.

- One of the first steps in processing a new set of data is *cleaning*.
- By cleaning the data, we ensure a few properties:
  - The data are in an analyzable format.
  - All data take legal values.
  - Any outliers are located and treated.
  - Any missing data are located and treated.



# A Little Notation


---

$Y :=$  An  $N \times P$  data matrix

$Y_{mis} :=$  The *missing* part of  $Y$

$Y_{obs} :=$  The *observed* part of  $Y$

$R :=$  An  $N \times P$  pattern matrix encoding nonresponse



# What are Missing Data?

---

Missing data are empty cells in a dataset where there should be observed values.

- The missing cells correspond to true population values, but we haven't observed those values.





# What are Missing Data?

---

Missing data are empty cells in a dataset where there should be observed values.

- The missing cells correspond to true population values, but we haven't observed those values.

Not every empty cell is a missing datum.

- Quality-of-life ratings for dead patients in a mortality study
- Firm profitability after the company goes out of business
- Self-reported severity of menstrual cramping for men
- Empty blocks of data following "gateway" items



# Missing Data Descriptives



# Missing Data Pattern

Missing data (or response) patterns represent unique combinations of observed and missing items.

- $P$  items  $\Rightarrow 2^P$  possible patterns.

	X	Y
1	x	y
2	x	.
3	.	y
4	.	.

Patterns for  $P = 2$

	X	Y	Z
1	x	y	z
2	x	y	.
3	x	.	z
4	.	y	z
5	x	.	.
6	.	.	z
7	.	y	.
8	.	.	.

Patterns for  $P = 3$

# Nonresponse Rates

---

## Percent/Proportion Missing

- The proportion of cells containing missing data
- Good early screening measure
- Should be computed for each variable, not for the entire dataset

## Attrition Rate

- The proportion of participants that drop-out of a study at each measurement occasion

## Percent/Proportion of Complete Cases

- The proportion of observations with no missing data
- Often reported but nearly useless quantity



# Nonresponse Rates


---

## Covariance Coverage

- The proportion of cases available to estimate a given pairwise relationship (e.g., a covariance between two variables)
- Very important to have adequate coverage for the parameters you want to estimate

## Fraction of Missing Information

- Associated with an estimated parameter, not with an incomplete variable
- Like an  $R^2$  for the missing data
- Most important diagnostic value for missing data problems
- Can only be computed after treating the missing data



# Covariance Coverage Examples

- What is the coverage for  $\text{cov}(X, Y)$ ?
- What is the coverage for  $\text{cov}(W, Y)$ ?
- What about  $\text{cov}(X, Z)$ ?

	W	X	Y	Z
1	w	x	y	.
2	w	x	y	.
3	w	x	y	.
4	w	x	y	.
5	w	x	y	.
6	w	.	y	z
7	w	.	y	z
8	w	.	y	z
9	w	.	y	z
10	w	.	y	z

# Nonresponse Rate Examples

- What is the percent missing at Time 2?
- What is the attrition rate at Time 3?

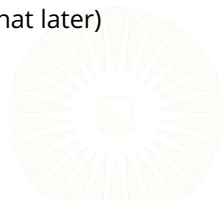
	T1	T2	T3	T4
1	x1	x2	x3	x4
2	x1	x2	x3	x4
3	x1	x2	x3	x4
4	x1	x2	x3	.
5	x1	x2	x3	.
6	x1	x2	.	.
7	x1	x2	.	.
8	x1	.	.	.
9	x1	.	.	.
10	x1	.	.	.

# What is an outlier?

---

For the time being, we're considering *univariate outliers*.

- Extreme values with respect to the distribution of a variable's other observations
  - A human height measurement of 3 meters
  - A high temperature in Utrecht of  $50^{\circ}$
  - Annual income of €250,000 for a student
- Not accounting for any particular model (we'll get to that later)





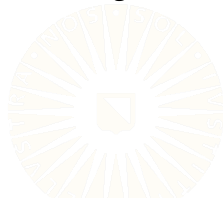
# What is an outlier?

---

A univariate outlier may, or may not, be an illegal value.

- Data entry errors are probably the most common cause.
- Outliers can also be legal, but extreme, values.

Key Point: We choose to view an outlier as arising from a different population than the one to which we want to generalize our findings.



# Finding Univariate Outliers

---

We have many methods available to diagnose potential outliers.

- Four of the simplest and most popular are:
  1. Internally studentized residuals (AKA Z-score method)
  2. Externally studentized residuals
  3. Median absolute deviation method
  4. Tukey's boxplot method



# Internally Studentized Residuals

---

For each observation,  $X_n$ , we compute the following quantity:

$$T_n = \frac{X_n - \bar{X}}{SD_X}$$

- $T_n$  follows a Student's  $t$  distribution with  $df = N - 1$ .
  - We can do a formal test for “outlier” status.
- Assuming a large sample, if  $T_n > C$  (where  $C$  is usually 2 or 3), we label  $X_n$  as an outlier.



# Internally Studentized Residuals

---

For each observation,  $X_n$ , we compute the following quantity:

$$T_n = \frac{X_n - \bar{X}}{SD_X}$$

- $T_n$  follows a Student's  $t$  distribution with  $df = N - 1$ .
  - We can do a formal test for “outlier” status.
- Assuming a large sample, if  $T_n > C$  (where  $C$  is usually 2 or 3), we label  $X_n$  as an outlier.

Although simple, this method has some substantial limitations.

- The cutpoint,  $C$ , can only be meaningfully chosen when  $X$  is normally distributed.
- Both  $\bar{X}$  and  $SD_X$  are highly sensitive to outliers.

# Externally Studentized Residual

The externally studentized residual method is essentially the same as the internally studentized residual method, but we adjust  $\bar{X}$  and  $SD_X$  to remove the influence of the observation we're evaluating.

- Let  $\mathbb{N}_{(n)} = \{1, \dots, (n-1), (n+1), \dots, N\}$ .
- Define the deletion mean,  $\bar{X}_{(n)}$ , and deletion SD,  $SD_{X(n)}$ , as:

$$\bar{X}_{(n)} = \frac{1}{N-1} \sum_{i \in \mathbb{N}_{(n)}} X_i$$
$$SD_{X(n)} = \sqrt{\frac{1}{N-2} \sum_{i \in \mathbb{N}_{(n)}} (X_i - \bar{X}_{(n)})^2}$$



# Externally Studentized Residual

---

The externally studentized residual is defined in the same way as the internally studentized version:

$$T_{(n)} = \frac{X_n - \bar{X}_{(n)}}{SD_{X(n)}}$$

- $T_{(n)}$  follows a Student's  $t$  distribution with  $df = N - 2$ .
  - We can do a formal test for “outlier” status.
- Assuming a large sample, if  $T_{(n)} > C$  (where  $C$  is usually 2 or 3), we label  $X_n$  as an outlier.



# Externally Studentized Residual


The externally studentized residual is defined in the same way as the internally studentized version:

$$T_{(n)} = \frac{X_n - \bar{X}_{(n)}}{SD_{X(n)}}$$

- $T_{(n)}$  follows a Student's  $t$  distribution with  $df = N - 2$ .
  - We can do a formal test for “outlier” status.
- Assuming a large sample, if  $T_{(n)} > C$  (where  $C$  is usually 2 or 3), we label  $X_n$  as an outlier.

$T_{(n)}$  is immune to the influence of the  $n$ th observation.

- Still requires  $X$  to be normally distributed
- Still sensitive to outliers other than the  $n$ th observation



# Median Absolute Deviation Method

---

The biggest limitation of studentized residuals is that their measures of central tendency and dispersion are sensitive to outliers.

- If we can replace the (deleted) mean and the (deleted) SD with more robust statistics, we can avoid this issue.
  - Replace the mean,  $\bar{X}$ , with the *median*,  $\text{Med}(X)$
  - Replace the SD with the *median absolute deviation*:

$$MAD_X = b \times \text{Med} (|X_n - \text{Med}(X)|)$$

- We choose the coefficient as  $b = 1/Q_{0.75}$
- For the normal distribution,  $b \approx 1/0.6745 \approx 1.4826$





# Median Absolute Deviation Method

---

We compute our test statistic by replacing the mean with the median and the SD with the MAD in the standard Wald test formula:

$$T_{MAD} = \frac{X_n - \text{Med}(X)}{MAD_X}$$

- $T_{MAD}$  doesn't allow for formal statistical tests.
- We can use the same general cutoffs we would use for the studentized residual methods.
  - Assuming a large sample, if  $T_{(n)} > C$  (where  $C$  is usually 2 or 3), we label  $X_n$  as an outlier.



# Median Absolute Deviation Method

---

We compute our test statistic by replacing the mean with the median and the SD with the MAD in the standard Wald test formula:

$$T_{MAD} = \frac{X_n - \text{Med}(X)}{MAD_X}$$

- $T_{MAD}$  doesn't allow for formal statistical tests.
- We can use the same general cutoffs we would use for the studentized residual methods.
  - Assuming a large sample, if  $T_{(n)} > C$  (where  $C$  is usually 2 or 3), we label  $X_n$  as an outlier.

$T_{MAD}$  is immune to the influence of, up to, 50% outlying observations.

- Requires us to assume a parametric distribution for  $X$ 
  - This assumption is necessary to compute  $b$ .

# Breakdown Point

---

To compare robust statistics, we consider their *breakdown points*.

- The breakdown point is the minimum proportion of cases that must be replaced by  $\infty$  to cause the value of the statistic to go to  $\infty$ .

The mean has a breakdown point of  $1/N$ .


- Replacing a single value with  $\infty$  will produce an infinite mean.

The deletion mean has a breakdown point of  $2/N$ .

- We can replace, at most, 1 value with  $\infty$  without producing an infinite mean.

The median has breakdown point of 50%.


- We can replace  $n < N/2$  of the observations with  $\infty$  without producing an infinite median.



# Boxplot Method

Tukey (1977) described a procedure for flagging potential outliers based on the familiar box-and-whiskers plot.

- Does not require normally distributed  $X$
- Not sensitive to outliers
- Doesn't allow for formal statistical tests



# Boxplot Method

---

A *fence* is an interval defined as the following function of the *first quartile*, the *third quartile*, and the *inner quartile range* ( $IQR = Q_3 - Q_1$ ):

$$F = \{Q_1 - C \times IQR, Q_3 + C \times IQR\}$$

- Taking  $C = 1.5$  produces the *inner fence*.
- Taking  $C = 3.0$  produces the *outer fence*.

We can use these fences to identify potential outliers:

- Any value that falls outside of the inner fence is a *possible outlier*.
- Any value that falls outside of the outer fence is a *probable outlier*.

# Multivariate Outliers

---

Sometimes, the combinations of values in an observation are very unlikely, even when no individual value is an outlier.

- These observations are *multivariate outliers*.
  - A person in the 95<sup>th</sup> percentile for height and the 5<sup>th</sup> percentile for weight
  - A person who simultaneously scores highly on scales of depression and positive affect

To detect multivariate outliers, we use *distance metrics*.

- Distance metrics quantify the similarity of two vectors.
  - Similarity between two observations
  - Similarity between an observation and the mean vector



# Mahalanobis Distance

---

One of the most common distance metrics is the *Mahalanobis Distance*.

- The Mahalanobis distance,  $\Delta$ , is a multivariate generalization of the internally studentized residual:

$$\Delta_n = \sqrt{(\mathbf{x}_n - \hat{\mu}_{\mathbf{X}})^T \hat{\Sigma}_{\mathbf{X}}^{-1} (\mathbf{x}_n - \hat{\mu}_{\mathbf{X}})}$$

As with studentized residuals, if  $\Delta_n > C$ , we label  $\mathbf{x}_n$  as an outlier.

- When  $\mathbf{X}$  is  $K$ -variate normally distributed,  $\Delta_n^2$  follows a  $\chi^2$  distribution with  $df = K$ .
- We take  $C$  to be the square-root of a suitably conservative quantile (e.g.,  $q \in \{99\%, 99.9\%\}$ ) of the  $\chi_K^2$  distribution:  $C = \sqrt{\chi_{K,q}^2}$ .

# Problems with Mahalanobis Distance

---

Like the internally studentized residual, Mahalanobis distance is highly sensitive to outliers.

- The underlying estimates of central tendency,  $\hat{\mu}_{\mathbf{X}}$ , and dispersion,  $\hat{\Sigma}_{\mathbf{X}}$ , are computed using all observations.





# Problems with Mahalanobis Distance

---

Like the internally studentized residual, Mahalanobis distance is highly sensitive to outliers.

- The underlying estimates of central tendency,  $\hat{\mu}_{\mathbf{X}}$ , and dispersion,  $\hat{\Sigma}_{\mathbf{X}}$ , are computed using all observations.

We want robust analogues of  $\hat{\mu}_{\mathbf{X}}$  and  $\hat{\Sigma}_{\mathbf{X}}$ .

- We have several options for robust estimation of  $\hat{\mu}_{\mathbf{X}}$  and  $\hat{\Sigma}_{\mathbf{X}}$ . E.g.:
  - Minimum covariance determinant method (MCD; Rousseeuw, 1985)
  - Minimum volume ellipsoid method (MVE; Rousseeuw, 1985)
  - M-estimation (Maronna, 1976)
- Conceptually, robust methods operate by either:
  - Using only a “good” subset of data to estimate  $\hat{\mu}_{\mathbf{X}}$  and  $\hat{\Sigma}_{\mathbf{X}}$ .
  - Downweighting outlying observations when estimating  $\hat{\mu}_{\mathbf{X}}$  and  $\hat{\Sigma}_{\mathbf{X}}$ .

# Robust Mahalanobis Distance


---

Equipped with robust estimates of central tendency,  $\hat{\mu}_{R,X}$ , and dispersion,  $\hat{\Sigma}_{R,X}$ , we define the robust Mahalanobis distance in the natural way:

$$\Delta_{R,n} = \sqrt{(\mathbf{x}_n - \hat{\mu}_{R,X})^T \hat{\Sigma}_{R,X}^{-1} (\mathbf{x}_n - \hat{\mu}_{R,X})}$$

We use  $\Delta_{R,n}$  in the same way as  $\Delta_n$ .

- If  $\Delta_{R,n} > C$ , we label  $\mathbf{x}_n$  as an outlier.
- Again, we take  $C$  to be the square-root of some quantile of the  $\chi_K^2$  distribution:  $C = \sqrt{\chi_{K,q}^2}$ .



# Practicalities: Univariate vs. Multivariate

---

Univariate outlier checks are safe for most variables.



# Practicalities: Univariate vs. Multivariate

---

Univariate outlier checks are safe for most variables.

Don't include too many variables in multivariate outlier checks.

- More variables increases the chances of false positives.
- E.g., don't run a multivariate outlier test on your entire dataset.



# Practicalities: Univariate vs. Multivariate

---

Univariate outlier checks are safe for most variables.

Don't include too many variables in multivariate outlier checks.

- More variables increases the chances of false positives.
- E.g., don't run a multivariate outlier test on your entire dataset.

Do use multivariate outlier checks for scales.

- E.g., if you have a psychometric scale measuring depression, you should check the items of that scale for multivariate outliers.



# Practicalities: Univariate vs. Multivariate

---

Univariate outlier checks are safe for most variables.

Don't include too many variables in multivariate outlier checks.

- More variables increases the chances of false positives.
- E.g., don't run a multivariate outlier test on your entire dataset.

Do use multivariate outlier checks for scales.

- E.g., if you have a psychometric scale measuring depression, you should check the items of that scale for multivariate outliers.

Maybe check the variables in a single model for multivariate outliers.

- E.g., if you have a small set of items that you will include in a regression model, it could make sense to check these variables for multivariate outliers.

# Practicalities: Outliers for Categorical Data

---

Nominal, ordinal, and binary items *can* have outliers.

- Outliers on categorical variables are often more indicative of bad variables than outlying cases.



# Practicalities: Outliers for Categorical Data

---

Nominal, ordinal, and binary items *can* have outliers.

- Outliers on categorical variables are often more indicative of bad variables than outlying cases.

Ordinal

- Most participant endorse one of the lowest categories on an ordinal item, but a few participants endorse the highest category.
- The participants who endorse the highest category may be outliers.





# Practicalities: Outliers for Categorical Data

---

Nominal, ordinal, and binary items *can* have outliers.

- Outliers on categorical variables are often more indicative of bad variables than outlying cases.

## Ordinal

- Most participant endorse one of the lowest categories on an ordinal item, but a few participants endorse the highest category.
- The participants who endorse the highest category may be outliers.

## Nominal

- Groups with very low membership may be outliers on nominal grouping variables.



# Practicalities: Outliers for Categorical Data

---

Nominal, ordinal, and binary items *can* have outliers.

- Outliers on categorical variables are often more indicative of bad variables than outlying cases.

## Ordinal

- Most participant endorse one of the lowest categories on an ordinal item, but a few participants endorse the highest category.
- The participants who endorse the highest category may be outliers.

## Nominal

- Groups with very low membership may be outliers on nominal grouping variables.

## Binary

- If most endorse the item, the few who do not may be outliers.

# Treating Outliers

---

If we locate any outliers, they must be treated.

- Outliers caused by errors, mistakes, or malfunctions (i.e., *error outliers*) should be directly corrected.
- Labeling non-error outliers is a subjective task.
  - A (non-error) outlier must originate from a population separate from the one we care about.
  - Don't blindly automate the decision process.



# Treating Outliers


---

If we locate any outliers, they must be treated.

- Outliers caused by errors, mistakes, or malfunctions (i.e., *error outliers*) should be directly corrected.
- Labeling non-error outliers is a subjective task.
  - A (non-error) outlier must originate from a population separate from the one we care about.
  - Don't blindly automate the decision process.

The most direct solution is to delete any outlying observation.

- If you delete non-error outliers, the analysis should be reported twice: with outliers and without.



# Treating Outliers

---

For univariate outliers, we can use less extreme types of deletion.


- Delete outlying values (but not the entire observation).
- These empty cells then become missing data.

Winsorization:

- Replace the missing values with the nearest non-outlying value.

Missing data analysis:

- Treat the missing values along with any naturally-occurring nonresponse.



# Treating Outliers

---

We can also use robust regression procedures to estimate the model directly in the presence of outliers.

- Weight the objective function to reduce the impact of outliers
  - M-estimation
- Trim outlying observations during estimation
  - Least trimmed squares, MCD, MVE
- Take the median, instead of the mean, of the squared residuals
  - Least median of squares
- Model some quantile of the DV's distribution instead of the mean
  - Quantile regression
- Model the outcome with a heavy-tailed distribution
  - Laplacian, Student's T



# References

---

- Maronna, R. A. (1976). Robust  $m$ -estimators of multivariate location and scatter. *The Annals of Statistics*, 4(1), 51–67. doi: 10.1214/aos/1176343347
- O'Neil, C., & Schutt, R. (2014). *Doing data science: Straight talk from the frontline*. Sebastopol, CA: O'Reilly Media, Inc.
- Rousseeuw, P. J. (1985). Multivariate estimation with high breakdown point. In W. Grossmann, G. Pflug, I. Vincze, & W. Wertz (Eds.), *Mathematical statistics and applications* (Vol. B, pp. 283–297). Netherlands: Reidel.
- Tukey, J. W. (1977). *Exploratory data analysis* (Vol. 2). Reading, MA: Addison-Wesley.